Symbolic Types for Lenient Symbolic Execution

ANONYMOUS AUTHOR(S)

We present $\widehat{\lambda}_{\bullet}$, a typed λ-calculus for *lenient symbolic execution*, where some language constructs do not recognize symbolic values. Its type system, however, ensures safe behavior of all symbolic values in a program. Our calculus extends a base occurrence typing system with symbolic types and mutable state, making it a suitable model for both functional and imperative symbolically executed languages. Naively allowing mutation in this mixed setting introduces soundness issues, however, so we further add *concreteness polymorphism*, which restores soundness without rejecting too many valid programs. To show that our calculus is a useful model for a real language, we implemented Typed Rosette, a typed extension of the solver-aided Rosette language. We evaluate Typed Rosette by porting a large code base, demonstrating that our type system accommodates a wide variety of symbolically executed programs.

1 DEBUGGING LENIENT SYMBOLIC EXECUTION

Programmers are increasingly using symbolic execution (Boyer et al. 1975; King 1975), in conjunction with satisfiability solvers, to help find bugs (Cadar et al. 2008; Farzan et al. 2013; Godefroid et al. 2012), verify program properties (Dennis et al. 2006; Jaffar et al. 2012; Near and Jackson 2012), and synthesize code from specifications (Solar-Lezama et al. 2005; Torlak and Bodik 2014). One challenge when designing such tools is deciding which language constructs should handle symbolic values. Lifting all language constructs is costly to implement and may produce complex solver encodings for some language features. Conversely, restricting symbolic execution to a language subset limits expressive power and the number of possible applications. Some testing frameworks allow a form of mixed symbolic execution where symbolic values are concretized before interacting with unlifted language constructs. This does not consider all program paths, however, and thus may not be suitable for all applications of symbolic execution.

Ideally, we would like *lenient symbolic execution*, where a small language subset is lifted for symbolic execution but those forms may freely interact with a larger unlifted language (Torlak and Bodik 2013). This approach is complete, is easier to implement, simplifies solver encodings, yet still allows programmers to use expressive language features. Debugging such mixed programs, however, can be tricky. Consider the following pseudocode example:

```
if (unlifted-int? x) then (add1 x) else (error "cannot add non-integer")
```

The author of this code might reasonably conclude that it will not error when `x` is an integer. Executing the code with `x` as a *symbolic* integer, however, is problematic if `unlifted-int?` does not recognize symbolic values, i.e., it returns true only when given a *concrete* integer and false otherwise. In this scenario, the program will unexpectedly reach the error case. Worse, when symbolic execution is paired with a solver, a programmer may only detect a problem (if at all) after examining the solver's output, and even then will not have much information to debug with.

Thus, despite the potential benefits, most symbolic execution engines have shunned the mixing of lifted and unlifted language constructs. In contrast, we propose *embracing lenient symbolic execution but supplementing it with a type system.* Such a system would enjoy the benefits of lenient symbolic execution, yet programmers would not be burdened with manually ensuring safe interaction of symbolic values. Further, any safety violations would be automatically detected and reported *before*

2017. 2475-1421/2017/1-ART1 $15.00
DOI: 10.1145/nnnnnnn.nnnnnnn

execution. Finally, such problems are easier to fix since since they would be reported in the context of their occurrence, rather than after constraint solving. Our paper makes two main contributions:

(1) $\widehat{\lambda}_{\bullet}$, a typed λ-calculus that allows but ensures safe lenient symbolic execution, and
(2) Typed Rosette, a solver-aided typed programming language based on our calculus.

$\widehat{\lambda}_{\bullet}$ includes symbolic values, higher-order functions, and mutation, and thus may represent the core of both functional and imperative symbolically executed languages. Its type system, in addition to ensuring safe interaction of symbolic values, utilizes occurrence typing and true union types: the path-sensitivity of the former fits well with the path-oriented nature of symbolic execution, and the latter naturally accommodates the symbolic union values that arise when computing with symbolic values. We further extend the system with *concreteness polymorphism*, which utilizes intersection types to add context sensitivity and more precisely tracks the flow of symbolic values.

To show that our type system is useful, we created Typed Rosette, a typed extension of the Rosette language (Torlak and Bodik 2014). Rosette utilizes symbolic execution to compile programs to solver constraints and programmers have used it it for a variety of verification and synthesis applications (Bornholt and Torlak 2017; Bornholt et al. 2016; Pernsteiner et al. 2016; Phothilimthana et al. 2014; Weitz et al. 2016). Rosette is an ideal target for our type system since:

- it equips only a subset of its host language, Racket, to handle symbolic values;
- it allows potentially unsafe interaction with the rest of Racket (Flatt and PLT 2010) but discourages programmers from doing so; and
- despite the warnings, users routinely stray from safe core in order to benefit from the extra unsafe features, relying on vigilant programming and detailed knowledge of the language to manually ensure only safe uses of symbolic values.

Typed Rosette aims to help with the last task. Our paper is organized as follows:

(1) section §2 motivates the paper with examples;
(2) section §3 presents our symbolic λ-calculus and its type system;
(3) section §4 presents its metatheory;
(4) section §5 introduces Typed Rosette;
(5) section §6 evaluates the usefulness of Typed Rosette;
(6) and finally, the paper concludes (§ 7-8) with a discussion of related and future work.

2 ROSETTE PRIMER AND MOTIVATING EXAMPLES

This section introduces lenient symbolic execution in unyped Rosette via examples and motivates the need for a type system. In Rosette, base symbolic values must be explicitly introduced and computing with these values may produce other symbolic values:

```
#lang rosette/safe
(define-symbolic i integer?)
(+ i 1) ; => i+1
(define-symbolic b boolean?)
(if b 1 (+ i 2)) ; => ⟨[b: 1][¬b: i + 2]⟩
```

This example uses the safe part of Rosette, declared with `#lang rosette/safe`, where all constructs properly handle symbolic values. For example, the third line, using the symbolic integer value defined on the second line, evaluates to the symbolic term `i+1`. Conditionals may also create symbolic values. For example, evaluation of the fifth line, due to its symbolic boolean test, explores both branches and merges the results into a guarded symbolic union value.[1]

[1]Rosette further distinguishes between "solvable" and general symbolic union values but we ignore this distinction for now.

The output of symbolic execution is a set of assertions that are passed to a solver. Programmers interact with the solver via a convenient API that allows various high-level queries, e.g., `solve`:

```
#lang rosette/safe
(define-symbolic i integer?)
(solve (assert (= i 3))) ; => model: i = 3
```

The `solve` query produces an assignment of values to symbolic variables that satisfies the assertions collected during symbolic execution, if one exists. In this example, the solver determines a value of 3 for `i`. Dually, the `verify` query tries to find a counterexample to the assertions, e.g.:

```
#lang rosette/safe
(define (sorted? v)
  (define-symbolic i j integer?)
  (define max (sub1 (vector-length v)))     ; largest index
  (implies (and (<= 0 i max) (<= 0 j max)   ; assume valid indices
                (< i j))
           (<= (vector-ref v i)             ; check if each pair is sorted
               (vector-ref v j))))
(verify (assert (sorted? (vector 3 5 4)))) ; => ✗: i = 1, j = 2
```

This code attempts to verify whether a vector is sorted using a user-defined `sorted?` predicate. The definition uses two symbolic integer values, `i` and `j`, to check whether each pair of elements in the input vector is sorted. In this case, `(vector 3 5 4)` is not sorted and the solver finds a counterexample when $i = 1$ and $j = 2$.

Another useful pattern is to store symbolic values themselves in data structures:

```
#lang rosette/safe
(define-symbolic x y z integer?)
(define vec (vector x y z))
(verify #:assume    (assert (and (< x y) (< y z)))
        #:guarantee (assert (sorted? vec))) ; ✓
```

This example creates a vector of three symbolic integers `x`, `y`, and `z`. Then, assuming one can deduce that $x < y$ and $y < z$ from the rest program, the attempt to verify sortedness (using the previously defined predicate) succeeds.

One might reasonably wish to write a similar program using hash tables:

```
#lang rosette  ; allows lenient symbolic execution
(define (sorted-hash? h)
  (define-symbolic i j integer?)
  (define max (sub1 (hash-count h)))        ; largest index
  (implies (and (<= 0 i max) (<= 0 j max)   ; assume valid indices
                (< i j))
           (<= (hash-ref h i)               ; check if each pair is sorted
               (hash-ref h j))))
(define-symbolic x y z integer?)
(define h (hash 0 x 1 y 2 z))
(verify (assert (sorted-hash? h)))          ; => ✗: x=-1, y=0, z=1, i=0, j=1
```

This example uses the larger `#lang rosette` language, which includes unlifted language constructs. In particular, despite the similarity of the code to the previous vector examples, the use of `hash-ref` is a mistake because it does not properly handle symbolic arguments. Therefore the verification does not succeed and instead, the solver returns a strange counterexample.

This last example illustrates both the benefits and drawbacks of lenient symbolic execution. On one hand, permitting the use of unlifted constructs from arbitrary libraries greatly expands the possible applications of Rosette. For instance, whereas the safe Rosette subset supports only a few basic data structures such as lists and vectors, the full Rosette language enables applying symbolic execution techniques to hash tables and any other library within the Racket ecosystem. On the other hand, as the documentation warns: "Rosette cannot, in general, guarantee the safety or correctness of programs that use unlifted constructs. As a result, the programmer is responsible for ensuring that these programs behave correctly." Worse, debugging can be challenging when things go wrong. In this case, the verification would have succeeded if the programmer had manually iterated over the hash table to produce the constraints, instead of using `hash-ref`. The strange solver output, however, provides no information to help the programmer arrive at this solution.

One might wonder why the untyped `hash-ref` function did not produce an evaluation-time exception to alert programmers to the issue. The answer is that there are two conflicting ways to utilize exceptions during symbolic evaluation. The first way assumes that exceptions represent "incorrect" programs and that they should be reported to the programmer. The second way, which Rosette chooses, assumes that exceptions are "correct" thus uses the information to prune out infeasible progam paths when generating constraints. These conflicting situations make evaluation-time checks insufficient for debugging symbolic programs.

To allow programmers the benefit of using arbitrary language features, while avoiding the burden of manually ensuring safety, we propose augmenting a language like Rosette with a type checker. With Typed Rosette, instead of a bad solver output, the compiler reports that `hash-ref` expected a concrete integer but got a symbolic integer with symbolic type $\widehat{\texttt{Int}}$ instead:

```
#lang typed/rosette
(define-symbolic i integer?)
  ; ...
(hash-ref h i) ; => TYERR: hash-ref expected Int, given i of type ^Int
```

3 A SYMBOLIC TYPED λ-CALCULUS

This section presents $\widehat{\lambda}_{\bullet}$, a typed λ-calculus whose types distinguish symbolic and concrete values. $\widehat{\lambda}_{\bullet}$ builds on the occurrence typing system of Tobin-Hochstadt and Felleisen (2010) (explained in § 3.1), whose path-sensitivity fits well with the path-based nature of symbolic execution. Further, occurrence typing easily accommodates true union types, which are suitable for the union values created during symbolic execution.

We first extend the base occurrence typed calculus with symbolic values and types that distinguish symbolic values from concrete ones (§ 3.2). Then, to improve the precision of our type system, we introduce the notion of *concreteness polymorphism* (§ 3.3). We further extend the functional core with mutation (§ 3.4), enabling our calculus to model both functional and imperative symbolic execution. Adding mutation naively, however, introduces unsoundness depending on the the concreteness of the path. Thus, finally, we extend concreteness polymorphism with additional path-sensitivity and context-sensitivity that fixes the possible unsoundness yet preserves enough precision so that mutation remains useful.

3.1 Occurrence Typing, in a Nutshell

Occurrence typing uses the notion of type refinement (Freeman and Pfenning 1991) to add more path-sensitivity than traditional type systems. Specifically, conditionals in an occurrence typing

$$
\begin{array}{rcll}
e \in Exp & ::= & i \mid s \mid \texttt{true} \mid \texttt{false} \mid x \mid op \mid \lambda x{:}\tau\,.\,e \mid e\,e \mid \texttt{if}\,e\,e\,e & \text{(terms)} \\
 & & i \in \mathbb{Z}, \quad s \in Strings & \\
op \in Op & ::= & \texttt{bool?} \mid \texttt{not} \mid \texttt{int?} \mid \texttt{add1} \mid \texttt{str?} \mid \texttt{strlen} & \text{(primitive ops)} \\
\tau \in Ty & ::= & \texttt{Int} \mid \texttt{String} \mid \texttt{True} \mid \texttt{False} \mid \tau_{fn} \mid \tau \cup \tau \mid \texttt{Any} & \text{(types)} \\
\tau_{fn} \in TyFn & ::= & x{:}\tau \xrightarrow{\psi\mid\psi;o} \tau & \text{(function types)} \\
\psi \in Prop & ::= & x{:}\tau \mid \neg x{:}\tau \mid \psi \supset \psi \mid \psi \vee \psi \mid \psi \wedge \psi \mid \top \mid \bot & \text{(propositions)} \\
\Gamma \in Env & ::= & \psi \ldots & \text{(environments)} \\
o \in Obj & ::= & x \mid \cdot & \text{(target objects)}
\end{array}
$$

Abbreviations:

$$\texttt{Bool} = \texttt{True} \cup \texttt{False}, \quad \texttt{Bot} = \cup$$

Fig. 1. Syntax of base occurrence typing system, based on Tobin-Hochstadt and Felleisen (2010).

system may refine the types of variables in its branches based on result of its test expression. Consider this function definition:[2]

```
(define (f [x : (U Int String)]) -> Int
  (if (int? x)
      (add1 x)
      (string-length x)))
```

Though the `x` input may initially be either a string or integer, its type is refined to `Int` in the "then" branch and `String` in the "else" branch, and thus both the addition and string length computation type check and the function returns an `Int`.

The conditional test expression dictates the type refinements in the branches. More precisely, type checking judgements have shape $\Gamma \vdash e \ : \ \tau;\ \psi^+ \mid \psi^-$ where, if e is used as a conditional test, then ψ^+ and ψ^- are logical propositions that describe how to refine types in the "then" and "else" branches, respectively. In the example, type checking `(int? x)` computes $\psi^+ = x{:}\texttt{Int}$ and $\psi^- = \neg x{:}\texttt{Int}$. When these constraints are combined with with `x`'s original type `(U Int String)`, the type checker may conclude $x{:}\texttt{Int}$ in the "then" branch and $x{:}\texttt{String}$ in the "else" branch.

Formally, figure 1 presents the syntax for a basic occurrence typing system; it more or less resembles the calculus of Tobin-Hochstadt and Felleisen (2010). The language includes three types of literal values: integers, strings, and booleans. Integers have base type `Int` and strings have base type `String`. To keep our type judgements uniform, `if` conditionals use "non-false" semantics, i.e., its test expression may have have any type and all non-false values are considered equivalent to true.[3] Consequently the type system includes two separate Boolean base types, `True` and `False`; we use an abbreviation `Bool` for the union of `True` and `False` when needed. The rest of the expressions are variables, primitive operations, lambdas, and function application; the rest of the types are function types, unions, and `Any`. Unions are a multi-arity type constructor, e.g., a union of no arguments is equivalent to the empty type, but we present it syntactically as a binary constructor when it's more convenient. Finally, `Any` includes all the other types.

[2]For consistency with the rest of the paper, we continue to use Typed Rosette's syntax for examples.
[3]The alternative requires separating boolean expression judgements, which include refinements, and non-boolean ones.

$$\boxed{\delta_\tau : Op \to TyFn}$$

$$\begin{aligned}
\delta_\tau\ \texttt{bool?} &= x{:}\texttt{Any} \xrightarrow{x{:}\texttt{Bool}\mid\neg x{:}\texttt{Bool}} \texttt{Bool} \\
\delta_\tau\ \texttt{int?} &= x{:}\texttt{Any} \xrightarrow{x{:}\texttt{Int}\mid\neg x{:}\texttt{Int}} \texttt{Bool} \\
\delta_\tau\ \texttt{str?} &= x{:}\texttt{Any} \xrightarrow{x{:}\texttt{String}\mid\neg x{:}\texttt{String}} \texttt{Bool} \\
\delta_\tau\ \texttt{not} &= x{:}\texttt{Any} \xrightarrow{x{:}\texttt{False}\mid\neg x{:}\texttt{False}} \texttt{Bool} \\
\delta_\tau\ \texttt{add1} &= x{:}\texttt{Int} \xrightarrow{\top\mid\bot} \texttt{Int} \\
\delta_\tau\ \texttt{strlen} &= x{:}\texttt{String} \xrightarrow{\top\mid\bot} \texttt{Int}
\end{aligned}$$

Fig. 2. Types for primitive ops in base occurrence typing system.

$\widehat{\lambda_o}$'s function type differs from the standard arrow type in three ways: it (1) binds a parameter that is in scope for the rest of the type, (2) specifies two refinement propositions ψ, and (3) specifies a "target object" o. The example from the beginning of the section introduced how type checking a conditional's test additionally computes propositions which are then used to refine a variable's type in the branches. The origin of these refinements are predicate functions such as `int?`, which has type $x{:}\texttt{Any} \xrightarrow{x{:}\texttt{Int}\mid\neg x{:}\texttt{Int}} \texttt{Bool}$. In other words applying `int?` reveals additional information about its argument: it either is an integer or is not an integer. A base type environment lookup function δ_τ, defined in figure 2, assigns types to all the primitive operations in a similar manner. The type of the `add1` and `strlen` functions specify $\top$ and $\bot$ propositions because their inputs, integers and strings, respectively, are always considered "true" when used in a conditional test, and the result of these primitive functions do not reveal additional information about the types of their inputs.

The final component of the function type is a target object o. Target objects, specified as the final component of a typing judgement $\Gamma \vdash e : \tau;\ \psi \mid \psi;\ o$, add flow-sensitivity to the type system in order to track the variable that "would be" modified by the ψ refinements. In the previous example, the result of `(int? x)` straightforwardly determines whether `x` is an integer or not. The target of refinement may be less straightforward, however, if `int?` is applied to a non-variable expression. For example if the conditional were changed to:

```
(if (int? (id x)) .... )
```

where `id` is the identify function, the type system should still know that `x` is the target of refinement, even though it passes through a function call.[4] To address these cases, function types also specify a target object. For example `id` would have type $x : \texttt{Any} \xrightarrow{\neg x{:}\texttt{False}\mid x{:}\texttt{False};\ x} \texttt{Any}$, where the object component x specifies that the function input becomes the target object after the function is applied. Together, the propositions and target objects make occurrence typing fully compositional. In other words, a conditional will properly refine the types in its branches with any arbitrary expression as its test expression.

[4]The "target object" component is particularly important when dealing with data structures, as originally presented by Tobin-Hochstadt and Felleisen (2010). We omit data structures to simplify presentation of our type system, since they are orthogonal to our goal of symbolic types, but they should work in the same manner described in the original calculus.

$$\textsc{T-Int}\quad \Gamma \vdash i \;:\; \texttt{Int};\; \top \mid \bot \qquad \textsc{T-True}\quad \Gamma \vdash \texttt{true} \;:\; \texttt{True};\; \top \mid \bot \qquad \textsc{T-False}\quad \Gamma \vdash \texttt{false} \;:\; \texttt{False};\; \bot \mid \top$$

$$\textsc{T-String}\quad \Gamma \vdash s \;:\; \texttt{String};\; \top \mid \bot \qquad \textsc{T-Prim}\quad \Gamma \vdash op \;:\; \delta_\tau(op);\; \top \mid \bot$$

$$\textsc{T-Var}\quad \frac{\Gamma \vdash_\psi x : \tau}{\Gamma \vdash x \;:\; \tau;\; \neg x{:}\texttt{False} \mid x{:}\texttt{False};\; x}$$

$$\textsc{T-If}\quad \frac{\begin{array}{c}\Gamma \vdash e_1 \;:\; \tau_1;\; \psi_1^+ \mid \psi_1^- \\ \Gamma, \psi_1^+ \vdash e_2 \;:\; \tau;\; \psi_2^+ \mid \psi_2^-;\; o \\ \Gamma, \psi_1^- \vdash e_3 \;:\; \tau;\; \psi_3^+ \mid \psi_3^-;\; o\end{array}}{\Gamma \vdash \texttt{if}\, e_1\, e_2\, e_3 \;:\; \tau;\; \psi_2^+ \vee \psi_3^+ \mid \psi_2^- \vee \psi_3^-;\; o}$$

$$\textsc{T-Lam}\quad \frac{\Gamma, x{:}\tau \vdash e \;:\; \tau';\; \psi^+ \mid \psi^-;\; o}{\Gamma \vdash \lambda x{:}\tau\,.\,e \;:\; x{:}\tau \xrightarrow{\psi^+ \mid \psi^-; o} \tau';\; \top \mid \bot}$$

$$\textsc{T-App}\quad \frac{\Gamma \vdash e_1 \;:\; x{:}\tau_1 \xrightarrow{\psi^+ \mid \psi^-; o} \tau_2;\; \psi_1^+ \mid \psi_1^- \qquad \Gamma \vdash e_2 \;:\; \tau_1;\; \psi_2^+ \mid \psi_2^-;\; o_2}{\Gamma \vdash e_1\, e_2 \;:\; \tau_2[x := o_2];\; \psi^+[x := o_2] \mid \psi^-[x := o_2];\; o[x := o_2]}$$

$$\textsc{T-Subsume}\quad \frac{\Gamma \vdash e \;:\; \tau;\; \psi^+ \mid \psi^-;\; o \qquad \Gamma, \psi^+ \vdash_\psi \psi^{+\prime} \qquad \Gamma, \psi^- \vdash_\psi \psi^{-\prime} \qquad \tau <: \tau' \qquad o <: o'}{\Gamma \vdash e \;:\; \tau';\; \psi^{+\prime} \mid \psi^{-\prime};\; o'}$$

Fig. 3. Type rules for base occurrence typing system.

Figure 3 presents the type rules for figure 1's base occurrence typing language.[5] The rules for literal integers, strings, and true values are straightforward: when these and other non-false values, e.g., functions, are used as a conditional test, the proposition environment is unchanged in "then" branch, denoted with proposition $\top$, and may be used to prove any conclusion in the "else" branch, denoted with $\bot$. Dually, the rule for false values flips these propositions.

Thus far we have only informally described propositions and how conditionals use them to refine types in their branches. Formally, a proposition environment Γ, which generalizes the type environment found in conventional type systems, propogates these propositions. The propositions are used to define the proof system in figure 4, and the T-Var rule in figure 3 uses this proof system to determine the type of a variable. The T-Var rule also states that if a variable is used as a conditional test, the only thing we can conclude in the "then" branch is that the variable is not false. Dually, the variable must be false in the "else" branch. Finally, a variable reference sets the current refinement target object to be that variable.

The rules in figure 4 are used to determine a variable's type and mostly consist of the standard rules of propositional logic. The additional L-Restrict and L-Remove rules combine various propositions to refine a more general type into a more precise one. The *restrict* and *remove* metafunctions, defined in figure 6, roughly correspond to set intersection and set difference operations, respectively. Revisiting the example from the beginning of the section, if the proposition environment contains $x : \texttt{Int} \cup \texttt{String}$ and $x : \texttt{Int}$, like in the "then" branch", then the L-Restrict rule allows the conclusion $x{:}\texttt{Int}$. If the proposition environment contains $x{:}\texttt{Int} \cup \texttt{String}$ and $\neg x{:}\texttt{Int}$, like in the "else" branch", then the L-Remove rule allows the conclusion $x{:}\texttt{String}$.

[5] We may omit showing the object component in function types, e.g., in figure 2, and type rules when they have the $\cdot$ object.

$$\textsc{L-Atom}\ \frac{\psi \in \Gamma}{\Gamma \vdash_\psi \psi} \qquad \textsc{L-True}\ \Gamma \vdash_\psi \top \qquad \textsc{L-False}\ \frac{\Gamma \vdash_\psi \bot}{\Gamma \vdash_\psi \psi}$$

$$\textsc{L-Sub}\ \frac{\Gamma \vdash_\psi x : \tau_1 \qquad \tau_1 <: \tau_2}{\Gamma \vdash_\psi x : \tau_2} \qquad \textsc{L-NotSub}\ \frac{\Gamma \vdash_\psi \neg x : \tau_2 \qquad \tau_1 <: \tau_2}{\Gamma \vdash_\psi \neg x : \tau_1}$$

$$\textsc{L-AndI}\ \frac{\Gamma \vdash_\psi \psi_1 \qquad \Gamma \vdash_\psi \psi_2}{\Gamma \vdash_\psi \psi_1 \wedge \psi_2} \qquad \textsc{L-AndE}\ \frac{\Gamma, \psi_1 \vdash_\psi \psi \text{ or } \Gamma, \psi_2 \vdash_\psi \psi}{\Gamma, \psi_1 \wedge \psi_2 \vdash_\psi \psi}$$

$$\textsc{L-OrI}\ \frac{\Gamma \vdash_\psi \psi_1 \text{ or } \Gamma \vdash_\psi \psi_2}{\Gamma \vdash_\psi \psi_1 \vee \psi_2} \qquad \textsc{L-OrE}\ \frac{\Gamma, \psi_1 \vdash_\psi \psi \qquad \Gamma, \psi_2 \vdash_\psi \psi}{\Gamma, \psi_1 \vee \psi_2 \vdash_\psi \psi}$$

$$\textsc{L-ImpI}\ \frac{\Gamma, \psi_1 \vdash_\psi \psi_2}{\Gamma \vdash_\psi \psi_1 \supset \psi_2} \qquad \textsc{L-ImpE}\ \frac{\Gamma \vdash_\psi \psi_1 \qquad \Gamma \vdash_\psi \psi_1 \supset \psi_2}{\Gamma \vdash_\psi \psi_2}$$

$$\textsc{L-Restrict}\ \frac{\Gamma \vdash_\psi x : \tau_1 \qquad \Gamma \vdash_\psi x : \tau_2}{\Gamma \vdash_\psi x : (\mathit{restrict}\ \tau_1\ \tau_2)} \qquad \textsc{L-Remove}\ \frac{\Gamma \vdash_\psi x : \tau_1 \qquad \Gamma \vdash_\psi \neg x : \tau_2}{\Gamma \vdash_\psi x : (\mathit{remove}\ \tau_1\ \tau_2)}$$

Fig. 4. Proof system for computing a variable's type in base occurrence typing system.

$$\textsc{Sub-Obj}\ o <: \cdot \qquad \textsc{Sub-Refl}\ \tau <: \tau \qquad \textsc{Sub-Any}\ \tau <: \mathsf{Any}$$

$$\textsc{Sub-}\cup\textsc{-L}\ \frac{\tau_1 <: \tau \qquad \tau_2 <: \tau}{\tau_1 \cup \tau_2 <: \tau} \qquad \textsc{Sub-}\cup\textsc{-R}\ \frac{\tau <: \tau_1 \text{ or } \tau <: \tau_2}{\tau <: \tau_1 \cup \tau_2}$$

$$\textsc{Sub-arr}\ \frac{\tau_3 <: \tau_1 \qquad \tau_2 <: \tau_4 \qquad \psi_1^+ \vdash \psi_2^+ \qquad \psi_1^- \vdash \psi_2^- \qquad o_1 <: o_2}{x : \tau_1 \xrightarrow{\psi_1^+ \mid \psi_1^- ; o_1} \tau_2 \ <: \ x : \tau_3 \xrightarrow{\psi_2^+ \mid \psi_2^- ; o_2} \tau_4}$$

Fig. 5. Subtype relation for base occurrence typing system.

$$\boxed{\mathit{restrict} : Ty\ Ty \to Ty}$$

$$\mathit{restrict}\ \tau_1\ \tau_2 = \mathsf{Bot}, \text{ if } \tau_1 \not<: \tau_2, \tau_2 \not<: \tau_1$$
$$\mathit{restrict}\ (\tau_1 \cup \tau_2)\ \tau = (\mathit{restrict}\ \tau_1\ \tau) \cup (\mathit{restrict}\ \tau_2\ \tau)$$
$$\mathit{restrict}\ \tau_1\ \tau_2 = \tau_1, \text{ if } \tau_1 <: \tau_2$$
$$\mathit{restrict}\ \tau_1\ \tau_2 = \tau_2, \text{ otherwise}$$

$$\boxed{\mathit{remove} : Ty\ Ty \to Ty}$$

$$\mathit{remove}\ \tau_1\ \tau_2 = \mathsf{Bot}, \text{ if } \tau_1 <: \tau_2$$
$$\mathit{remove}\ (\tau_1 \cup \tau_2)\ \tau = (\mathit{remove}\ \tau_1\ \tau) \cup (\mathit{remove}\ \tau_2\ \tau)$$
$$\mathit{remove}\ \tau_1\ \tau_2 = \tau_1, \text{ otherwise}$$

Fig. 6. Metafunctions for base occurrence typing system.

Returning to figure 3, the T-If rule shows how the positive and negative propositions from the test expression are propagated to the branches of a conditional, as previously described. The propositions for the conditional expression itself then, are disjunctions of the propositions from each branch. The T-Lam rule shows how the latent propositions and target object from a lambda's body are transferred to its function type. Dually, T-App specifies that applying a function uses the propositions and object from the function type as the propositions and object of the application expression, except that the object from the argument replaces the function parameter.

Finally, our base occurrence typing system uses a subtyping relation, defined in figure 5 and used in figure 3's T-Subsume rule. Notably, the Sub-∪-R shows how unions may be introduced, to complement the union-elimination nature of the `if` form. Also, the Sub-Obj rule overloads the subtyping relation <: to accommodate o objects and states that $\cdot$ is a "super type" of every other object. The remaing subtyping rules are more or less standard.

$$e \in Exp ::= \ldots \mid \widehat{x^{\tau}} \qquad \text{(terms)}$$
$$\tau \in Ty ::= \ldots \mid \widehat{\tau} \qquad \text{(types)}$$

$$\boxed{concrete? : Ty \to \texttt{Bool}}$$

$$\begin{aligned} concrete?\,\widehat{\tau} &= \texttt{false} \\ concrete?\,\tau_1 \cup \tau_2 &= concrete?\,\tau_1 \text{ and } concrete?\,\tau_2 \\ concrete?\,\tau_1 \cap \tau_2 &= concrete?\,\tau_1 \text{ or } concrete?\,\tau_2 \\ concrete?\,\tau &= \texttt{true}, \text{otherwise} \end{aligned}$$

Fig. 7. (left) $\widehat{\lambda_{\bullet}}$, lambda-calculus with symbolic values, extends core occurrence typing calculus from figure 1; (right) Metafunction to determine whether a type represents a concrete value.

T-SymInt
$$\Gamma \vdash \widehat{x^{\texttt{Int}}} \ : \ \widehat{\texttt{Int}};\ \top \mid \bot$$

T-SymBool
$$\Gamma \vdash \widehat{x^{\texttt{Bool}}} \ : \ \widehat{\texttt{Bool}};\ \top \mid \top$$

T-If-Sym
$$\frac{\Gamma \vdash e_1 \ : \ \widehat{\tau_1};\ \psi_1^+ \mid \psi_1^- \qquad \texttt{False} <: \widehat{\tau_1} \qquad \Gamma, \psi_1^+ \vdash e_2 \ : \ \tau;\ \psi_2^+ \mid \psi_2^-;\ o \qquad \Gamma, \psi_1^- \vdash e_3 \ : \ \tau;\ \psi_3^+ \mid \psi_3^-;\ o}{\Gamma \vdash \texttt{if}\, e_1\, e_2\, e_3 \ : \ \widehat{\tau};\ \psi_2^+ \vee \psi_3^+ \mid \psi_2^- \vee \psi_3^-;\ o}$$

T-If-Conc
$$\frac{\Gamma \vdash e_1 \ : \ \tau_1;\ \psi_1^+ \mid \psi_1^- \qquad concrete?\,\tau_1 \text{ or } (symbolic?\,\tau_1 \text{ and } \texttt{False} \not<: \tau_1) \qquad \Gamma, \psi_1^+ \vdash e_2 \ : \ \tau;\ \psi_2^+ \mid \psi_2^-;\ o \qquad \Gamma, \psi_1^- \vdash e_3 \ : \ \tau;\ \psi_3^+ \mid \psi_3^-;\ o}{\Gamma \vdash \texttt{if}\, e_1\, e_2\, e_3 \ : \ \tau;\ \psi_2^+ \vee \psi_3^+ \mid \psi_2^- \vee \psi_3^-;\ o}$$

Sub-ConcAny
$$\frac{concrete?\,\tau}{\tau <: \texttt{Any}}$$

Sub-SymAny
$$\tau <: \widehat{\texttt{Any}}$$

Sub-Sym
$$\tau <: \widehat{\tau}$$

Fig. 8. Some type and subtyping rules for $\widehat{\lambda_{\bullet}}$, a λ-calculus with symbolic values

3.2 Adding Symbolic Values

We introduce $\widehat{\lambda_{\bullet}}$ by extending the occurrence typing calculus in figure 1 with symbolic values and symbolic types. Specifically, the syntax for $\widehat{\lambda_{\bullet}}$ in figure 7 (left) extends figure 1 with symbolic literal values $\widehat{x^{\tau}}$. Programmers specify the kind of symbolic value they wish to introduce into the program with a type annotation. At the type level, $\widehat{\lambda_{\bullet}}$ distinguishes symbolic values from concrete ones with a $\widehat{\cdot}$ constructor and thus a $\widehat{x^{\tau}}$ value has type $\widehat{\tau}$. $\widehat{\lambda_{\bullet}}$ allows only integer and boolean base symbolic values, as enforced by the T-SymInt and T-SymBool rules in figure 8. Consequently, strings are always concrete and string functions in $\widehat{\lambda_{\bullet}}$ represent the "unlifted constructs" in a language supporting lenient symbolic execution.

The other form that may introduce symbolic values is `if` as specified by T-If-Sym. Specifically, if the test expression is symbolic *and* possibly a `False` value, then, since symbolic execution must explore both paths, the `if` expression is assigned a symbolic type using the $\widehat{\cdot}$ constructor. The $\texttt{False} <: \widehat{\tau_1}$ side-condition might seem counterintuitive at first glance. Instead, one might expect something like $\widehat{\tau_1} <: \widehat{\texttt{Bool}}$ instead. This latter side condition is too restrictive, however, for conditionals with non-false semantics. For example it does not include test expressions like $\texttt{if}\, \widehat{x^{\texttt{Bool}}}\, \widehat{x^{\texttt{Bool}}}\, \widehat{x^{\texttt{Int}}}$. This expression has type $\widehat{\texttt{Bool}} \cup \widehat{\texttt{Int}}$, which is not a subtype of $\widehat{\texttt{Bool}}$, yet symbolic execution must still explore both conditional branches. Thus the $\texttt{False} <: \widehat{\tau_1}$ side-condition is appropriate because it indicates when symbolic execution must explore both conditional branches.

$$\tau \in Ty ::= \ldots \mid \tau_{fn} \cap \tau_{fn} \qquad op \in Op ::= \ldots \mid \mathsf{conc?} \qquad \text{(types and prim ops)}$$

$$\text{T-Lam-ConcArg}\ \frac{\Gamma, x{:}\tau \vdash e\ :\ \tau';\ \psi^+ \mid \psi^-;\ o}{\Gamma \vdash \lambda x{:}\tau.e\ :\ x{:}\tau \xrightarrow{\psi^+|\psi^-;o} \tau';\ \top \mid \bot} \qquad \text{T-Lam-SymArg}\ \frac{\Gamma, x{:}\widehat{\tau} \vdash e\ :\ \tau';\ \psi^+ \mid \psi^-;\ o}{\Gamma \vdash \lambda x{:}\tau.e\ :\ x{:}\widehat{\tau} \xrightarrow{\psi^+|\psi^-;o} \tau';\ \top \mid \bot}$$

$$\text{T-Inter-I}\ \frac{\Gamma \vdash e\ :\ \tau_{fn_1} \qquad \Gamma \vdash e\ :\ \tau_{fn_2}}{\Gamma \vdash e\ :\ \tau_{fn_1} \cap \tau_{fn_2}} \qquad \text{Sub-}\cap\text{-L}\ \frac{\tau_{fn_1} <:\ \tau_{fn} \text{ or } \tau_{fn_2} <:\ \tau_{fn}}{\tau_{fn_1} \cap \tau_{fn_2} <:\ \tau_{fn}} \qquad \text{Sub-}\cap\text{-R}\ \frac{\tau_{fn} <:\ \tau_{fn_1} \qquad \tau_{fn} <:\ \tau_{fn_2}}{\tau_{fn} <:\ \tau_{fn_1} \cap \tau_{fn_2}}$$

Fig. 9. $\widehat{\lambda_{\bullet}}$ with concreteness polymorphism, part 1.

If a conditional test expression is a non-`false` symbolic value or a concrete value, symbolic execution need only explore one branch and thus T-If-Conc does not apply the $\widehat{\cdot}$ constructor to the type of the `if` expression. T-If-Sym and T-If-Conc, which replace T-If from figure 3, use a *concrete?* metafunction on types, defined in figure 7 (right), which returns true if its argument type represents concrete values. We also use the abbreviation *symbolic?*τ to mean "not *concrete?*τ".

The $\widehat{\cdot}$ type actually denotes *possibly* symbolic values. In other words, a concrete value may have a $\widehat{\cdot}$ type, which may occur when symbolic execution utilizes dynamic information unavailable during type checking. For example, if both branches of a symbolic conditional evaluate to the same concrete value, the result of evaluation should be that concrete value. The type system, however, must conservatively label the expression as possibly symbolic. The Sub-Sym subtype rule in figure 8 specifies this relation. Specifically a type τ is considered a subtype of a possibly symbolic version of that type $\widehat{\tau}$. In addition, $\widehat{\lambda_{\bullet}}$ requires two separate subtype rules to handle the `Any` type: the `Any` type represents only concrete values while $\widehat{\texttt{Any}}$ includes symbolic values.

3.3 Concreteness Polymorphism, Part 1

Adding symbolic values and unlifted constructs that do not handle symbolic values enables $\widehat{\lambda_{\bullet}}$ to model lenient symbolic execution. Adding symbolic *types* allows the type system to ensure the safe behavior of symbolic values by preventing them from reaching unlifted constructs, which should be assigned concrete types. To fully benefit from lenient symbolic execution, however, unlifted constructs should be allowed to interact with *concrete* values. To achieve this, the type system should try to preserve concreteness at the type level as much as possible.

Naively type checking symbolic values, however, quickly causes the entire program to become symbolic, which then limits the usefulness of lenient symbolic execution. As an example, what type should $\widehat{\lambda_{\bullet}}$ assign to `add1`? It should accommodate symbolic values and thus one might naively assign type $\widehat{\texttt{Int}} \to \widehat{\texttt{Int}}$. With this type, however, the result of applying `add1` to a *concrete* value would have a symbolic type and thus could not be used with an unlifted arithmetic function even though it is perfectly safe. To improve the precision of $\widehat{\lambda_{\bullet}}$'s type system, i.e., to preserve concreteness as much as possible, we use intersection function types, shown in figure 9. We call this kind of finitary polymorphism *concreteness polymorphism*.

Concreteness polymorphism enables more precise types for primitive operations as shown in figure 10. In contrast with figure 2, figure 10 assigns types that allows boolean and integer primitive functions to handle symbolic values. These lifted operations with intersection types may be applied at any of their constituent function types, as specified by the Sub-$\cap$-L rule in figure 9 (in conjunction with T-Subsume and T-App). As a result the type checker safely allows the result of `add1` to be

$$\boxed{\delta_\tau : Op \to Ty}$$

$$\delta_\tau\texttt{bool?} = x{:}\texttt{Any} \xrightarrow{x:\texttt{Bool}\mid\neg x:\texttt{Bool}} \texttt{Bool} \cap x{:}\widehat{\texttt{Any}} \xrightarrow{x:\widehat{\texttt{Bool}}\mid\neg x:\widehat{\texttt{Bool}}} \widehat{\texttt{Bool}}$$

$$\delta_\tau\texttt{int?} = x{:}\texttt{Any} \xrightarrow{x:\texttt{Int}\mid\neg x:\texttt{Int}} \texttt{Bool} \cap x{:}\widehat{\texttt{Any}} \xrightarrow{x:\widehat{\texttt{Int}}\mid\neg x:\widehat{\texttt{Int}}} \widehat{\texttt{Bool}}$$

$$\delta_\tau\texttt{str?} = x{:}\texttt{Any} \xrightarrow{x:\texttt{String}\mid\neg x:\texttt{String}} \texttt{Bool}$$

$$\delta_\tau\texttt{not} = x{:}\texttt{Any} \xrightarrow{x:\texttt{False}\mid\neg x:\texttt{False}} \texttt{Bool} \cap x{:}\widehat{\texttt{Any}} \xrightarrow{x:\widehat{\texttt{False}}\mid\neg x:\widehat{\texttt{False}}} \widehat{\texttt{Bool}}$$

$$\delta_\tau\texttt{add1} = x{:}\texttt{Int} \xrightarrow{\top\mid\bot} \texttt{Int} \cap x{:}\widehat{\texttt{Int}} \xrightarrow{\top\mid\bot} \widehat{\texttt{Int}}$$

$$\delta_\tau\texttt{strlen} = x{:}\texttt{String} \xrightarrow{\top\mid\bot} \texttt{Int}$$

$$\delta_\tau\texttt{conc?} = x{:}\texttt{Any} \xrightarrow{\top\mid\bot} \texttt{True} \cap x{:}\widehat{\texttt{Any}} \xrightarrow{x:\texttt{Any}\mid\neg x:\texttt{Any}} \widehat{\texttt{Bool}}$$

Fig. 10. Types for primitive ops in $\widehat{\lambda_{\bullet}}$.

passed to an unlifted arithmetic operation if the original argument was concrete. Primitive string functions remain unlifted in $\widehat{\lambda_{\bullet}}$ and thus do not accept symbolic values. Thus applying either `str?` or `strlen` to a symbolic value results in a type error. Finally, $\widehat{\lambda_{\bullet}}$ adds a `conc?` primitive which, in conjunction with `if`, serves as the elimination rule for values with possibly-symbolic types.

Lambdas support concreteness polymorphism as well. Specifically, lambda bodies are type checked twice: once with a concrete input and once with a symbolic input, as specified by T-Lam-ConcArg and T-Lam-SymArg, respectively, in figure 9 (these rules replace T-Lam in figure 3).[6] The T-Inter-I rule allows assigning a lambda an intersection type consisting of these two types. (Note that a separate T-Inter-I rule is necessary since Sub-∩-R is insufficient for introducing the intersection type in this case.)

3.4 Adding Mutation

To allow $\widehat{\lambda_{\bullet}}$ to serve as a model for imperative symbolically executed languages as well, we next add mutation, specifically `set!` and sequencing expressions, and a `Unit` type, as shown in figure 11. A naive combination of symbolic paths and mutation, however, is unsound. For example, a standard type rule for `set!` might look like:

$$\textsc{Mut-Wrong}\quad \frac{x{:}\tau \in \Gamma \qquad \Gamma \vdash e \;:\; \tau}{\Gamma \vdash \texttt{set!}\; x\; e \;:\; \texttt{Unit}}$$

This may not be safe, as seen in the following example:

```
(define-symbolic b boolean?)
(define x 0)            ; x is a concrete value with concrete type Int
(if b (set! x 10) (set! x 11))
x  ; => ⟨[b: 10][¬b: 11]⟩  ; x is a symbolic value, but still has concrete type Int
```

[6] Of course, to avoid exponential explosion of function type sizes, a practical language would allow programmers to more precisely choose which combinations of parameter concreteness should be included in the type.

$$e \in Exp ::= \ldots \mid () \mid \texttt{set!}\, x\, e \mid e\,;e \qquad \tau \in Ty ::= \ldots \mid \texttt{Unit} \qquad \text{(terms and types)}$$

Fig. 11. $\widehat{\lambda_{\bullet}}$ symbolic lambda-calculus, with mutation

$$\tau_{fn} ::= \pi; x{:}\tau \xrightarrow{\psi \mid \psi} \tau \quad \text{(function types)}$$
$$\pi \in Path ::= \bullet \mid \circ \quad \text{(path condition)}$$

SUB-PATH
$\bullet <: \circ$

Fig. 12. (left) $\widehat{\lambda_{\bullet}}$ concreteness polymorphic function types; (right) subtype relation for paths

The example mutates x under a symbolic path, changing it from a concrete to a symbolic value. The type remains concrete, however, and thus the operation is unsound. Observe that x becomes bound to a symbolic value, even though it is only ever assigned concrete values. One way to restore safety is to force all mutable variables to have symbolic types but this conflicts with our goal of preserving concreteness as much as possible for lenient symbolic execution.

Instead, our type system tracks the concreteness of the path and `set!` rule looks something like:

$$\text{MUT-OK}\quad \frac{x{:}\tau \in \Gamma \qquad \Gamma \vdash e \;:\; \tau \qquad \textit{symbolic?}\,\tau \text{ OR path is concrete}}{\Gamma \vdash \texttt{set!}\, x\, e \;:\; \texttt{Unit}}$$

In other words, we introduce extra path sensitivity to our type system to improve its precision.

This path sensitivity requires additional context sensitivity in function calls. For example, the following `set!` in the λ body may be safe or unsafe depending on its call site:

```
(define x 0)  ; x is a concrete value with concrete type Int
(define (f [y : Int]) (set! x y))
(f 1)
x  ; SAFE: x still concrete with concrete type
(define-symbolic b : Bool)
(if b (f 2) (f 3)))
x  ; => ⟨[b: 2][¬b: 3]⟩ (UNSAFE: x is a symbolic val but has concrete type)
```

The type system should ideally allow defining `f` since it may be used safely, but disallow calls to `f` in unsafe contexts. In other words, the type system should reject the expressions `(f 2)` and `(f 3)` above. To achieve this, we extend our notion of concreteness polymorphism to include the concreteness of the path.

3.5 Concreteness Polymorphism, Part 2

Figure 12 shows an extended function type that includes a path concreteness marker where $\bullet$ means concrete path and $\circ$ means symbolic path. More precisely, a function with type marked with $\bullet$ may only be applied in the context of a concrete path whereas a function with type marked with $\circ$ may be applied in any context.

Consequently, we split $\widehat{\lambda_{\bullet}}$'s type rules into two sets of judgements, with one set using a $\vdash^{\bullet}$ relation for type checking under a concrete path, and one set using a $\vdash^{\circ}$ relation for type checking under a symbolic path. Most of the $\vdash^{\bullet}$ concrete path rules and $\vdash^{\circ}$ symbolic path rules are identical to their

$$\textsc{T-Lam-ConcArgPath}^{\bullet}\ \frac{\Gamma, x{:}\tau \vdash^{\bullet} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\bullet} \lambda x{:}\tau\,.\,e \ :\ \bullet; x{:}\tau \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'} \qquad \textsc{T-Lam-ConcArg-SymPath}^{\bullet}\ \frac{\Gamma, x{:}\tau \vdash^{\circ} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\bullet} \lambda x{:}\tau\,.\,e \ :\ \circ; x{:}\tau \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'}$$

$$\textsc{T-Lam-SymArg-ConcPath}^{\bullet}\ \frac{\Gamma, x{:}\widehat{\tau} \vdash^{\bullet} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\bullet} \lambda x{:}\tau\,.\,e \ :\ \bullet; x{:}\widehat{\tau} \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'} \qquad \textsc{T-Lam-SymArgPath}^{\bullet}\ \frac{\Gamma, x{:}\widehat{\tau} \vdash^{\bullet} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\bullet} \lambda x{:}\tau\,.\,e \ :\ \circ; x{:}\widehat{\tau} \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'}$$

$$\textsc{T-Lam-ConcArgPath}^{\circ}\ \frac{\Gamma, x{:}\tau \vdash^{\bullet} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\circ} \lambda x{:}\tau\,.\,e \ :\ \bullet; x{:}\tau \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'} \qquad \textsc{T-Lam-ConcArg-SymPath}^{\circ}\ \frac{\Gamma, x{:}\tau \vdash^{\circ} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\circ} \lambda x{:}\tau\,.\,e \ :\ \circ; x{:}\tau \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'}$$

$$\textsc{T-Lam-SymArg-ConcPath}^{\circ}\ \frac{\Gamma, x{:}\widehat{\tau} \vdash^{\bullet} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\circ} \lambda x{:}\tau\,.\,e \ :\ \bullet; x{:}\widehat{\tau} \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'} \qquad \textsc{T-Lam-SymArgPath}^{\circ}\ \frac{\Gamma, x{:}\widehat{\tau} \vdash^{\circ} e \ :\ \tau';\ \psi^{+} \mid \psi^{-}}{\Gamma \vdash^{\circ} \lambda x{:}\tau\,.\,e \ :\ \circ; x{:}\widehat{\tau} \xrightarrow{\psi^{+}\mid\psi^{-}} \tau'}$$

$$\textsc{T-App}^{\bullet}\ \frac{\Gamma \vdash^{\bullet} e_1 \ :\ \bullet; x{:}\tau_1 \xrightarrow{\psi^{+}\mid\psi^{-}} \tau_2 \qquad \Gamma \vdash^{\bullet} e_2 \ :\ \tau_1}{\Gamma \vdash^{\bullet} e_1\, e_2 \ :\ \tau_2;\ \psi^{+} \mid \psi^{-}} \qquad \textsc{T-App}^{\circ}\ \frac{\Gamma \vdash^{\circ} e_1 \ :\ \circ; x{:}\tau_1 \xrightarrow{\psi^{+}\mid\psi^{-}} \tau_2 \qquad \Gamma \vdash^{\circ} e_2 \ :\ \tau_1}{\Gamma \vdash^{\circ} e_1\, e_2 \ :\ \tau_2;\ \psi^{+} \mid \psi^{-}}$$

$$\textsc{T-If-Conc}^{\bullet}\ \frac{\begin{array}{c}\Gamma \vdash^{\bullet} e_1 \ :\ \tau_1;\ \psi_1^{+} \mid \psi_1^{-} \\ \textit{concrete?}\,\tau_1 \text{ or } (\textit{symbolic?}\,\tau_1 \text{ and } \texttt{False} \not<: \tau_1) \\ \Gamma, \psi_1^{+} \vdash^{\bullet} e_2 \ :\ \tau;\ \psi_2^{+} \mid \psi_2^{-} \\ \Gamma, \psi_1^{-} \vdash^{\bullet} e_3 \ :\ \tau;\ \psi_3^{+} \mid \psi_3^{-}\end{array}}{\Gamma \vdash^{\bullet} \texttt{if}\, e_1\, e_2\, e_3 \ :\ \tau;\ \psi_2^{+} \vee \psi_3^{+} \mid \psi_2^{-} \vee \psi_3^{-}} \qquad \textsc{T-If-Sym}^{\bullet}\ \frac{\begin{array}{c}\Gamma \vdash^{\bullet} e_1 \ :\ \widehat{\tau}_1;\ \psi_1^{+} \mid \psi_1^{-} \qquad \texttt{False} <: \widehat{\tau}_1 \\ \Gamma, \psi_1^{+} \vdash^{\bullet} e_2 \ :\ \tau;\ \psi_2^{+} \mid \psi_2^{-} \\ \Gamma, \psi_1^{-} \vdash^{\bullet} e_3 \ :\ \tau;\ \psi_3^{+} \mid \psi_3^{-}\end{array}}{\Gamma \vdash^{\bullet} \texttt{if}\, e_1\, e_2\, e_3 \ :\ \widehat{\tau};\ \psi_2^{+} \vee \psi_3^{+} \mid \psi_2^{-} \vee \psi_3^{-}}$$

$$\textsc{T-If-Conc}^{\circ}\ \frac{\begin{array}{c}\Gamma \vdash^{\circ} e_1 \ :\ \tau_1;\ \psi_1^{+} \mid \psi_1^{-} \\ \textit{concrete?}\,\tau_1 \text{ or } (\textit{symbolic?}\,\tau_1 \text{ and } \texttt{False} \not<: \tau_1) \\ \Gamma, \psi_1^{+} \vdash^{\circ} e_2 \ :\ \tau;\ \psi_2^{+} \mid \psi_2^{-} \\ \Gamma, \psi_1^{-} \vdash^{\circ} e_3 \ :\ \tau;\ \psi_3^{+} \mid \psi_3^{-}\end{array}}{\Gamma \vdash^{\circ} \texttt{if}\, e_1\, e_2\, e_3 \ :\ \tau;\ \psi_2^{+} \vee \psi_3^{+} \mid \psi_2^{-} \vee \psi_3^{-}} \qquad \textsc{T-If-Sym}^{\circ}\ \frac{\begin{array}{c}\Gamma \vdash^{\circ} e_1 \ :\ \widehat{\tau}_1;\ \psi_1^{+} \mid \psi_1^{-} \qquad \texttt{False} <: \widehat{\tau}_1 \\ \Gamma, \psi_1^{+} \vdash^{\circ} e_2 \ :\ \tau;\ \psi_2^{+} \mid \psi_2^{-} \\ \Gamma, \psi_1^{-} \vdash^{\circ} e_3 \ :\ \tau;\ \psi_3^{+} \mid \psi_3^{-}\end{array}}{\Gamma \vdash^{\circ} \texttt{if}\, e_1\, e_2\, e_3 \ :\ \widehat{\tau};\ \psi_2^{+} \vee \psi_3^{+} \mid \psi_2^{-} \vee \psi_3^{-}}$$

Fig. 13. Type rules for $\widehat{\lambda}_{\mathbf{o}}$ that consider path concreteness ($\vdash^{\bullet}$: concrete path, $\vdash^{\circ}$: symbolic path).

counterparts from figures 3, 8, and 9 except for the lambda, function application, and conditional rules. Figure 13 shows new versions these rules.[7]

Whereas the T-Lam-ConcArg and T-Lam-SymArg rules in figure 9 assign lambdas two possible arrow types, one for a symbolic input and one for a concrete one, the T-Lam-ConcArgPath$^{\bullet}$, T-Lam-ConcArg-SymPath$^{\bullet}$, T-Lam-SymArg-ConcPath$^{\bullet}$, and T-Lam-SymArgPath$^{\bullet}$ rules in figure 13 assign lambdas four possible arrow types, where the concreteness of the path is now considered in addition to the concreteness of the input. Specifically, a lambda body is type checked in four different contexts, one corresponding to each kind of input and path combination. Similarly, the T-Lam-ConcArgPath$^{\circ}$, T-Lam-ConcArg-SymPath$^{\circ}$, T-Lam-SymArg-ConcPath$^{\circ}$, and T-Lam-SymArgPath$^{\circ}$ rules assign types to lambdas defined under a symbolic path. Finally, the

[7] To simplify the presentation, figure 13 omits the o component of its judgements rules since its behavior remains unchanged.

$$\text{T-Set!}^{\bullet}\ \frac{x{:}\tau \in \Gamma \qquad \Gamma \vdash^{\bullet} e \ : \ \tau}{\Gamma \vdash^{\bullet} \texttt{set!}\ x\ e \ : \ \texttt{Unit};\ \top \mid \bot}$$

$$\text{T-Set!}^{\circ}\ \frac{x{:}\widehat{\tau} \in \Gamma \qquad \Gamma \vdash^{\circ} e \ : \ \widehat{\tau}}{\Gamma \vdash^{\circ} \texttt{set!}\ x\ e \ : \ \texttt{Unit};\ \top \mid \bot}$$

$$\text{T-Seq}^{\bullet}\ \frac{\Gamma \vdash^{\bullet} e_1 \ : \ \tau_1 \qquad \Gamma \vdash^{\bullet} e_2 \ : \ \tau_2;\ \psi_2^+ \mid \psi_2^-}{\Gamma \vdash^{\bullet} e_1\,;e_2 \ : \ \tau_2;\ \psi_2^+ \mid \psi_2^-}$$

$$\text{T-Seq}^{\circ}\ \frac{\Gamma \vdash^{\circ} e_1 \ : \ \tau_1 \qquad \Gamma \vdash^{\circ} e_2 \ : \ \tau_2;\ \psi_2^+ \mid \psi_2^-}{\Gamma \vdash^{\circ} e_1\,;e_2 \ : \ \tau_2;\ \psi_2^+ \mid \psi_2^-}$$

Fig. 14. $\widehat{\lambda}_{\bullet}$ type rules for mutation, considering path concreteness.

T-App$^{\bullet}$ and T-App$^{\circ}$ rules enforce that functions are only applied in appropriate contexts, with the former requiring $\bullet$-marked functions and the latter allowing $\circ$-marked functions.

The last four rules in figure 13 are the conditional rules, T-If-Conc$^{\bullet}$ and T-If-Sym$^{\bullet}$ for concrete paths, and T-If-Conc$^{\circ}$ and T-If-Sym$^{\circ}$ for symbolic paths. The rules propagate the path concreteness of the entire conditional expression to their branches, except for T-If-Sym$^{\bullet}$, which switches from $\vdash^{\bullet}$ to $\vdash^{\circ}$ when type checking the branches to properly handle the symbolic path.

With the rules in figure 13, we may define sound mutation rules, shown in figure 14. T-Set!$^{\bullet}$ specifies that any mutation is acceptable in a concrete path while T-Set!$^{\circ}$ allows mutation of only variables with symbolic type. Observe that these rules do not use type judgements in the premise to determine the acceptable type for the new value e, since subtyping would allow lifting the type to be symbolic. Instead, the proposition environment must directly include a proposition that x may have the specified type. Finally, T-Seq$^{\bullet}$ and T-Seq$^{\circ}$ specify that a sequence expression has the type and propositions of its last expression, in both concrete and symbolic paths.

4 METATHEORY

This section describes some desirable properties of $\widehat{\lambda}_{\bullet}$. The key theorem is a soundness result demonstrating that symbolic values cannot cause evaluation to get stuck by flowing to positions that cannot handle them.

4.1 Dynamic Semantics

Figure 15 extends the previous grammars with the notion of values, both symbolic and concrete, which are the result of symbolic execution. Specifically, a value v is either a concrete integer, string, boolean, void value, lambda, primitive function, or a symbolic value $\widehat{v}$. The symbolic values are the base symbolic booleans and integers; symbolic expressions, which result from evaluating primitive operations applied to symbolic values; or guarded symbolic union values, which result from evaluating symbolic conditionals.

Figure 17 shows how programs evaluate to values. Specifically, using the syntax in figure 16, it defines a CESK-style (Felleisen et al. 2009), register-machine semantics[8] for $\widehat{\lambda}_{\bullet}$. A machine state is a 4-tuple consisting of a (C)ontrol (or (C)urrent) expression e, a value (E)nvironment ρ whose domain includes free variables in e, a (S)tore σ, and a stac(K) κ. Briefly, value environments map variables to store locations, stores map locations to closures c, closures are a value-environment pair, and stack frames represent the context of evaluation.

[8]Our semantics is more or less equivalent to that of Torlak and Bodik (2014) with three key differences: (1) our merging function $\mu_{\widehat{v}}$ is simpler since it is not our main focus; (2) we exclude the solver API and do not include an explicit path register; for our purposes, it is sufficient for individual symbolic values to track their own guard conditions; and (3) we include an explicit value environment, which helps define machine state type judgements in order to show soundness.

$$v \in Val ::= i \mid s \mid \texttt{true} \mid \texttt{false} \mid () \mid op \mid \lambda x{:}\tau\,.\,e \mid \widehat{v} \qquad \text{(values)}$$
$$\widehat{v} \in SVal ::= \widehat{x}^{\texttt{Bool}} \mid \widehat{x}^{\texttt{Int}} \mid \widehat{op}\,\widehat{v} \mid \langle[\widehat{v}\texttt{:}\ v]\ldots\rangle \qquad \text{(symbolic values)}$$

Fig. 15. $\widehat{\lambda_{\bullet}}$ values

$$M \in Mach ::= \langle e, \rho, \sigma, \kappa\rangle \qquad V \in VMach ::= \langle v, \rho, \sigma, \langle\rangle\rangle \qquad \text{(machine states, final states)}$$
$$\rho \in VEnv ::= x \mapsto \ell, \ldots \quad \ell \in Loc \quad \sigma \in Sto ::= \ell \mapsto c, \ldots \quad c \in Clo ::= \langle v, \rho\rangle \qquad \text{(envs, stores)}$$
$$\kappa \in Frames ::= \langle\rangle \mid \left\langle \widehat{\texttt{if}_1}\texttt{:}\ \widehat{v}, \rho_{\widehat{v}}, e, \rho, \sigma, \kappa \right\rangle^{\circ} \mid \left\langle \widehat{\texttt{if}_2}\texttt{:}\ \widehat{v}, \rho_{\widehat{v}}, v, \rho, \sigma, \kappa \right\rangle^{\circ} \mid \qquad \text{(stack frames)}$$
$$\langle \texttt{if}\texttt{:}\ e, e, \rho, \kappa\rangle^{\pi} \mid \langle \texttt{arg}\texttt{:}\ e, \rho, \kappa\rangle^{\pi} \mid \langle \texttt{fn}\texttt{:}\ v, \rho, \kappa\rangle^{\pi} \mid \langle \texttt{set!}\texttt{:}\ \ell, \kappa\rangle^{\pi} \mid \langle \texttt{seq}\texttt{:}\ e, \rho, \kappa\rangle^{\pi}$$

Fig. 16. Grammar for CESK machine semantics for $\widehat{\lambda_{\bullet}}$.

The machine syntax is mostly standard, except for the stack frames which are decorated with a path concreteness marker, analogous to the type rules in figure 13. The marker determines the concreteness of the path for the subexpressions in the stack frame. These markers do not affect evaluation of programs, however, i.e., no left-hand sides in figure 17 use this path concreteness information. Instead, they are included only to help with the soundness proof. For this reason, figure 17 omits the path markers; instead, a new stack frame implicitly inherits the same marker as its preceding frame, except for $\widehat{\texttt{if}}$ frames, which always have a symbolic path marker.

The rules in figure 17 are also mostly straightforward: evaluation of subexpressions proceeds in call-by-value order, using the stack to save its context. For example App-Fn begins evaluation of an application expression by setting the control expression to be the function and saves the argument expression and a copy of the environment in a new `arg` stack frame.

Figure 17's rules deviate from traditional CESK rules where evaluation mixes both concrete and symbolic values. For example, evaluation of conditionals may produce symbolic values, as illustrated by the $\widehat{\text{If}}$, $\widehat{\text{If}}$-Then, and $\widehat{\text{If}}$-Else rules. Specifically, when the test expression is a symbolic boolean $\widehat{v}$, *both* branches are evaluated and the resulting values, guarded by $\widehat{v}$, make up the branches of the resulting symbolic union value. Evaluation of each branch should occur independently of the other branch, i.e., they should evaluate with different σ stores. Thus the $\widehat{\texttt{if}_1}$ and $\widehat{\texttt{if}_2}$ stack frames save a store σ to accompany the "else" and "then" branches, respectively.

After both branches are evaluated, the $\widehat{\text{if}}$-else rule uses the $\mu_{\widehat{v}}$ merging function, defined in figure 18, to create a guarded symbolic union value. The function creates the symbolic value from the test value and branch results, taking special care to merge their environments to avoid name conflicts. In addition, $\widehat{\text{if}}$-else merges the stores from each branch using the μ_{σ} function. The function creates guarded symbolic union values for any locations that point to different values.

One other key rule is App-Op; it specifies how primitive operations are evaluated using a δ metafunction, whose definition is presented in figure 19. The δ function evaluates application of primitive operations to *concrete* values in the expected manner. In addition, it dictates which and how primitive operations should handle *symbolic* values. Some operations, e.g., `strlen` and `str?`, do not support symbolic values and evaluation gets stuck if these operations are applied to symbolic values. In other cases, such as applying `bool?` to $\widehat{x}^{\texttt{Bool}}$,a symbolic input may turn into a concrete value result. More likely, however, applying a primitive operation to a symbolic value results in another symbolic value. For example, applying `add1` to $\widehat{x}^{\texttt{Int}}$ results in the symbolic

$$\langle x, \rho, \sigma, \kappa\rangle \mapsto \langle v, \rho_v, \sigma, \kappa\rangle \quad \text{where } \sigma(\rho(x)) = \langle v, \rho_v\rangle \qquad \text{(VAR)}$$

$$\langle e_1\, e_2, \rho, \sigma, \kappa\rangle \mapsto \langle e_1, \rho, \sigma, \langle \texttt{arg:}\ e_2, \rho, \kappa\rangle\rangle \qquad \text{(APP-FN)}$$

$$\langle v, \rho_v, \sigma, \langle \texttt{arg:}\ e, \rho, \kappa\rangle\rangle \mapsto \langle e, \rho, \sigma, \langle \texttt{fn:}\ v, \rho_v, \kappa\rangle\rangle \qquad \text{(APP-ARG)}$$

$$\left\langle v, \rho, \sigma, \left\langle \texttt{fn:}\ op, \rho_{op}, \kappa\right\rangle\right\rangle \mapsto \langle v', \rho, \sigma, \kappa\rangle \quad \text{where } \delta\, op\, v = v' \qquad \text{(APP-OP)}$$

$$\langle v, \rho_v, \sigma, \langle \texttt{fn:}\ \lambda x\!:\!\tau\,.\,e, \rho, \kappa\rangle\rangle \mapsto \langle e, \rho[x \mapsto \ell], \sigma[\ell \mapsto \langle v, \rho_v\rangle], \kappa\rangle \quad \text{where } \ell \notin dom(\sigma) \qquad \text{(APP-}\beta\text{)}$$

$$\langle \texttt{if}\ e_1\, e_2\, e_3, \rho, \sigma, \kappa\rangle \mapsto \langle e_1, \rho, \sigma, \langle \texttt{if:}\ e_2, e_3, \rho, \kappa\rangle\rangle \qquad \text{(IF)}$$

$$\langle \texttt{false}, \rho_v, \sigma, \langle \texttt{if:}\ e_{then}, e_{else}, \rho, \kappa\rangle\rangle \mapsto \langle e_{else}, \rho, \sigma, \kappa\rangle \qquad \text{(IF-FALSE)}$$

$$\langle v, \rho_v, \sigma, \langle \texttt{if:}\ e_{then}, e_{else}, \rho, \kappa\rangle\rangle \mapsto \langle e_{then}, \rho, \sigma, \kappa\rangle \quad \text{where } v \neq \texttt{false},\ v \neq \widehat{v} \qquad \text{(IF-TRUE1)}$$

$$\langle \widehat{v}, \rho_v, \sigma, \langle \texttt{if:}\ e_{then}, e_{else}, \rho, \kappa\rangle\rangle \mapsto \langle e_{then}, \rho, \sigma, \kappa\rangle \quad \text{where } \texttt{bool?}\widehat{v} = \texttt{false} \qquad \text{(IF-TRUE2)}$$

$$\langle \widehat{v}, \rho, \sigma, \langle \texttt{if:}\ e_1, e_2, \rho_1, \kappa\rangle\rangle \mapsto \left\langle e_1, \rho_1, \sigma, \left\langle \widehat{\texttt{if}}_1\texttt{:}\ \widehat{v}, \rho, e_2, \rho_1, \sigma, \kappa\right\rangle\right\rangle \qquad (\widehat{\text{IF}})$$

$$\left\langle v, \rho, \sigma, \left\langle \widehat{\texttt{if}}_1\texttt{:}\ \widehat{v}, \rho_{\widehat{v}}, e_2, \rho_2, \sigma_2, \kappa\right\rangle\right\rangle \mapsto \left\langle e_2, \rho_2, \sigma_2, \left\langle \widehat{\texttt{if}}_2\texttt{:}\ \widehat{v}, \rho_{\widehat{v}}, v, \rho, \sigma, \kappa\right\rangle\right\rangle \qquad (\widehat{\text{IF}}\text{-THEN})$$

$$\left\langle v_2, \rho_2, \sigma_2, \left\langle \widehat{\texttt{if}}_2\texttt{:}\ \widehat{v}, \rho, v_1, \rho_1, \sigma_1, \kappa\right\rangle\right\rangle \mapsto \langle \widehat{v}_3, \rho_3, \mu_\sigma(\langle \widehat{v}, \rho\rangle, \sigma_1, \sigma_2), \kappa\rangle \quad \langle \widehat{v}_3, \rho_3\rangle = \mu_{\widehat{v}}(\langle \widehat{v}, \rho\rangle, \langle v_1, \rho_1\rangle, \langle v_2, \rho_2\rangle) \qquad (\widehat{\text{IF}}\text{-ELSE})$$

$$\langle \texttt{set!}\ x\, e, \rho, \sigma, \kappa\rangle \mapsto \langle e, \rho, \sigma, \langle \texttt{set!:}\ \rho(x), \kappa\rangle\rangle \qquad \text{(SET!)}$$

$$\langle v, \rho, \sigma, \langle \texttt{set!:}\ \ell, \kappa\rangle\rangle \mapsto \langle (), \rho, \sigma[\ell \mapsto \langle v, \rho\rangle], \kappa\rangle \qquad \text{(SET-IT!)}$$

$$\langle e_1\,;e_2, \rho, \sigma, \kappa\rangle \mapsto \langle e_1, \rho, \sigma, \langle \texttt{seq:}\ e_2, \rho, \kappa\rangle\rangle \qquad \text{(SEQ1)}$$

$$\langle v, \rho_v, \sigma, \langle \texttt{seq:}\ e, \rho, \kappa\rangle\rangle \mapsto \langle e, \rho, \sigma, \kappa\rangle \qquad \text{(SEQ2)}$$

Fig. 17. $\widehat{\lambda}_{\bullet}$ CESK machine semantics

$$\boxed{\theta : Clo, VEnv, VEnv \to Clo}$$

$$\theta(\langle v, \rho\rangle, \rho_1, \rho_2) = \langle v[x := y] \ldots, \{y \mapsto \rho[x] \mid x \in dom(\rho), y \notin dom(\rho_1, \rho_2)\}\rangle$$

$$\boxed{\mu_{\widehat{v}} : Clo, Clo, Clo \to Clo}$$

$$\mu_{\widehat{v}}(\langle \widehat{v}, \rho\rangle, \langle v_1, \rho_1\rangle, \langle v_2, \rho_2\rangle) = \left\langle \langle [\widehat{v}_3 \colon v_4][\widehat{\mathsf{not}}\ \widehat{v}_3 \colon v_5]\rangle, \rho_3 \cup \rho_4 \cup \rho_5 \right\rangle$$
$$\text{where } \langle \widehat{v}_3, \rho_3\rangle = \theta(\langle \widehat{v}, \rho\rangle, \rho_1, \rho_2)$$
$$\langle v_4, \rho_4\rangle = \theta(\langle v_1, \rho_1\rangle, \rho, \rho_2)$$
$$\langle v_5, \rho_5\rangle = \theta(\langle v_2, \rho_2\rangle, \rho, \rho_1)$$

$$\boxed{\mu_\sigma : Clo, Sto, Sto \to Sto}$$

$$\begin{aligned}
\mu_\sigma(\langle \widehat{v}, \rho\rangle, \sigma_1, \sigma_2) = {} & \{\ell \mapsto \sigma_1[\ell] \mid \ell \in dom(\sigma_1), \ell \notin dom(\sigma_2)\} \cup \\
& \{\ell \mapsto \sigma_2[\ell] \mid \ell \in dom(\sigma_2), \ell \notin dom(\sigma_1)\} \cup \\
& \{\ell \mapsto \sigma_1[\ell] \mid \ell \in dom(\rho_1), \ell \in dom(\rho_2), \sigma_1[\ell] = \sigma_2[\ell]\} \cup \\
& \{\ell \mapsto \mu_{\widehat{v}}(\langle \widehat{v}, \rho\rangle, \sigma_1[\ell], \sigma_2[\ell]) \mid \ell \in dom(\rho_1), \ell \in dom(\rho_2), \sigma_1[\ell] \neq \sigma_2[\ell]\}
\end{aligned}$$

Fig. 18. Merge function for environments and stores.

$$\begin{aligned}
\delta\,\mathsf{conc?}\,\widehat{v} &= \mathtt{false} & \delta\,\mathsf{conc?}\,v &= \mathtt{true}, \text{ where } v \neq \widehat{v} \\
\delta\,\mathsf{bool?}\,\mathtt{true} &= \mathtt{true} & \delta\,\mathsf{str?}\,s &= \mathtt{true} \\
\delta\,\mathsf{bool?}\,\mathtt{false} &= \mathtt{true} & \delta\,\mathsf{str?}\,v &= \mathtt{false}, \text{ where } v \neq s, v \neq \widehat{v} \\
\delta\,\mathsf{bool?}\,v &= \mathtt{false} & \delta\,\mathsf{strlen}\,s &= \#\text{ chars in } s \\
\text{where } v \neq \mathtt{true},\ v \neq \mathtt{false}, v \neq \widehat{v} & & \delta\,\mathsf{int?}\,i &= \mathtt{true} \\
\delta\,\mathsf{bool?}\,\widehat{x}^{\mathtt{Bool}} &= \mathtt{true} & \delta\,\mathsf{int?}\,v &= \mathtt{false}, v \neq i, v \neq \widehat{v} \\
\delta\,\mathsf{bool?}\,\widehat{x}^{\tau} &= \mathtt{false}, \text{ where } \tau \neq \mathtt{Bool} & \delta\,\mathsf{int?}\,\widehat{x}^{\mathtt{Int}} &= \mathtt{true} \\
\delta\,\mathsf{bool?}\,(\widehat{\mathsf{not}}\ \widehat{v}) &= \mathtt{true} & \delta\,\mathsf{int?}\,\widehat{x}^{\tau} &= \mathtt{false}, \text{ where } \tau \neq \mathtt{Int} \\
\delta\,\mathsf{bool?}\,(\widehat{op}\ \widehat{v}) &= \mathtt{false}, \text{ where } op \neq \mathtt{not} & \delta\,\mathsf{int?}\,(\widehat{\mathsf{add1}}\ \widehat{v}) &= \mathtt{true} \\
\delta\,\mathsf{bool?}\,\langle[\widehat{v}\colon v]\ldots\rangle &= \langle[\widehat{v}\colon \delta\,\mathsf{bool?}\,v]\ldots\rangle & \delta\,\mathsf{int?}\,(\widehat{op}\ \widehat{v}) &= \mathtt{false}, \text{ where } op \neq \mathtt{add1} \\
\delta\,\mathtt{not}\,\mathtt{false} &= \mathtt{true} & \delta\,\mathsf{int?}\,\langle[\widehat{v}\colon v]\ldots\rangle &= \langle[\widehat{v}\colon \delta\,\mathsf{int?}\,v]\ldots\rangle \\
\delta\,\mathtt{not}\,v &= \mathtt{false}, v \neq \widehat{v}, v \neq \mathtt{false} & \delta\,\mathtt{add1}\,i &= i + 1 \\
\delta\,\mathtt{not}\,\widehat{x}^{\mathtt{Bool}} &= \widehat{\mathtt{not}}\,\widehat{x}^{\mathtt{Bool}} & \delta\,\mathtt{add1}\,\widehat{x}^{\mathtt{Int}} &= \widehat{\mathtt{add1}}\,\widehat{x}^{\mathtt{Int}} \\
\delta\,\mathtt{not}\,\widehat{x}^{\tau} &= \mathtt{true}, \text{ where } \tau \neq \mathtt{Bool} & \delta\,\mathtt{add1}\,\langle[\widehat{v}\colon v]\ldots\rangle &= \langle[\widehat{v}\colon \delta\,\mathtt{add1}\,v']\ldots\rangle \\
\delta\,\mathtt{not}\,\langle[\widehat{v}\colon v]\ldots\rangle &= \langle[\widehat{v}\colon \delta\,\mathtt{not}\,v]\ldots\rangle & \text{where } v' \in v\ldots, \mathsf{int?}\,v'
\end{aligned}$$

Fig. 19. Evaluation of $\widehat{\lambda}_{\bullet}$ primitive operations.

$$\text{T-Op-SymVal}\quad \frac{\Gamma \vdash^{\pi} op\,\widehat{v} \;:\; \tau;\; \psi^{+} \mid \psi^{-};\; o}{\Gamma \vdash^{\pi} \widehat{op}\,\widehat{v} \;:\; \tau;\; \psi^{+} \mid \psi^{-};\; o}$$

$$\text{T-Union-SymVal}\quad \frac{\Gamma \vdash^{\pi} \widehat{v}_1 \;:\; \tau_1;\; \psi_1^{+} \mid \psi_1^{-};\; o_1 \qquad \Gamma, \psi_1^{+} \vdash^{\circ} v_1 \;:\; \tau;\; \psi_3^{+} \mid \psi_3^{-};\; o \qquad \Gamma \vdash^{\pi} \widehat{v}_2 \;:\; \tau_2;\; \psi_2^{+} \mid \psi_2^{-};\; o_2 \qquad \Gamma, \psi_2^{+} \vdash^{\circ} v_2 \;:\; \tau;\; \psi_4^{+} \mid \psi_4^{-};\; o}{\Gamma \vdash^{\pi} \langle [\widehat{v}_1 \colon v_1][\widehat{v}_2 \colon v_2] \rangle \;:\; \tau;\; \psi_3^{+} \vee \psi_3^{-} \mid \psi_4^{+} \vee \psi_4^{-};\; o}$$

$$\text{TM-CESK}\quad \frac{\mathcal{E}(\rho, \sigma) \vdash^{\pi} e : \tau_2;\; \psi_2^{+} \mid \psi_2^{-};\; o_2 \qquad \tau_2;\; \psi_2^{+} \mid \psi_2^{-};\; o_2; \sigma \vdash_{\kappa} \kappa^{\pi} : \tau;\; \psi^{+} \mid \psi^{-};\; o}{\vdash_{M} \langle e, \rho, \sigma, \kappa^{\pi} \rangle : \tau;\; \psi^{+} \mid \psi^{-};\; o}$$

Fig. 20. Type judgements for machine states and symbolic values.

$$\boxed{\mathcal{E} : VEnv, Sto \to Env}$$

$$\mathcal{E}(\rho, \sigma) = \{x : \tau \mid x \in dom(\rho)\}$$
$$\text{where } \langle v, \rho_2 \rangle = \sigma[\rho[x]] \text{ and } \mathcal{E}(\rho_2, \sigma) \vdash v \;:\; \tau$$

Fig. 21. Metafunction converting value and environment and store to proposition environment.

value $\widehat{\texttt{add1}}\,\widehat{x}^{\texttt{Int}}$. Finally, applying a primitive operation to a guarded union value traverses the tree, recursively applying the primitive operation. The most interesting case is applying `add1` to such a guarded union value. In this case, the tree is additionally pruned of non-integer branches.

4.2 Soundness

To evaluate a program e using the semantics in figure 17, an *inject* function first compiles the program to initial machine configuration $\langle e, \cdot, \cdot, \langle\rangle\rangle$ with an empty environment, store, and stack. Using this function, theorem 4.1 states the main theoretical result, which is that evaluating a well-typed program cannot get stuck.

Theorem 4.1 (Soundness). *If* $\vdash e \;:\; \tau;\; \psi^{+} \mid \psi^{-};\; o$, $M = inject(e)$, *and evaluating* M *terminates, then* $M \mapsto^{*} V$, *and* $\vdash_{M} V : \tau;\; \psi^{+\prime} \mid \psi^{-\prime};\; o'$ *for some* $\psi^{+\prime}$, $\psi^{-\prime}$, *and* o'.

We prove theorem 4.1 using a standard progress and preservation approach (Wright and Felleisen 1994). To do so, we first need to define typing judgements for symbolic values and machine states, both shown in figure 20. The T-Op-SymVal uses the rule for application from figure 13 and the T-Union-SymVal mostly mirrors the type rule for `if`.

The $\vdash_{M}$ relation for machine states is, essentially, a "bottom-up" version of the traditional "top-down" type judgements for expressions from the previous sections, e.g., figure 13. Specifically, $\vdash_{M}$ first uses those expression type judgements to type check the control expression. It then feeds the output of that judgement to a $\vdash_{\kappa}$ judgement, defined in figures 22 and 23, for type checking the stack. Intuitively, this input type information fills the "hole" that is implicit in each stack frame.

The $\vdash_{M}$ and $\vdash_{\kappa}$ relations both rely on a $\mathcal{E}$ metafunction, defined in figure 21, that "uncompiles" a runtime value environment and store into a proposition environment for type checking. This $\mathcal{E}$ metafunction computes types for elements of the value environment using the type judgements from figure 13. This function does not have access to path concreteness information, however, but as lemma 4.2 states, it does not matter which variant of the type judgement is used.

Lemma 4.2 (Types for Values). *For all* v, $\Gamma \vdash^{\bullet} v \;:\; \tau;\; \psi^{+} \mid \psi^{-};\; o$ *iff* $\Gamma \vdash^{\circ} v \;:\; \tau;\; \psi^{+} \mid \psi^{-};\; o$.

$$\text{TK-Empty}\quad \tau;\ \psi^+ \mid \psi^-;\ o;\ \sigma \vdash_\kappa \langle\rangle : \tau;\ \psi^+ \mid \psi^-;\ o$$

$$\text{TK-If-Conc}^\pi\quad \frac{\begin{array}{c}\textit{concrete?}\,\tau_{in} \text{ or } (\textit{symbolic?}\,\tau_{in} \text{ and } \texttt{False} \not<: \tau_{in}) \\ \mathcal{E}(\rho,\sigma), \psi^+_{in} \Vdash^\pi e_1 \ :\ \tau;\ \psi^+_1 \mid \psi^-_1;\ o \\ \mathcal{E}(\rho,\sigma), \psi^-_{in} \Vdash^\pi e_2 \ :\ \tau;\ \psi^+_2 \mid \psi^-_2;\ o \\ \tau;\ \psi^+_1 \vee \psi^+_2 \mid \psi^-_1 \vee \psi^-_2\ o;\ \sigma \vdash_\kappa \kappa : \tau';\ \psi^+ \mid \psi^-;\ o'\end{array}}{\tau_{in};\ \psi^+_{in} \mid \psi^-_{in};\ o_{in};\ \sigma \vdash_\kappa \langle \texttt{if}\colon e_1, e_2, \rho, \kappa\rangle^\pi : \tau';\ \psi^+ \mid \psi^-;\ o'}$$

$$\text{TK-If-Sym}^\pi\quad \frac{\begin{array}{c}\texttt{False} <: \widehat{\tau}_{in} \\ \mathcal{E}(\rho,\sigma), \psi^+_{in} \Vdash^\circ e_1 \ :\ \tau;\ \psi^+_1 \mid \psi^-_1; o \\ \mathcal{E}(\rho,\sigma), \psi^-_{in} \Vdash^\circ e_2 \ :\ \tau;\ \psi^+_2 \mid \psi^-_2; o \\ \widehat{\tau};\ \psi^+_1 \vee \psi^+_2 \mid \psi^-_1 \vee \psi^-_2\ o;\ \sigma \vdash_\kappa \kappa : \tau';\ \psi^+ \mid \psi^-;\ o'\end{array}}{\widehat{\tau}_{in}; \psi^+_{in} \mid \psi^-_{in}; o_{in}; \sigma \vdash_\kappa \langle \texttt{if}\colon e_1, e_2, \rho, \kappa\rangle^\pi : \tau';\ \psi^+ \mid \psi^-;\ o'}$$

$$\text{TK-SymIf1}^\circ\quad \frac{\begin{array}{c}\mathcal{E}(\rho_2,\sigma_2) \Vdash^\circ e_2 \ :\ \tau_1;\ \psi^+_2 \mid \psi^-_2; o_1 \\ \widehat{\tau}_1; \psi^+_1 \vee \psi^+_2 \mid \psi^-_1 \vee \psi^-_2; o_1; \mu_\sigma(\langle \widehat{v}, \rho\rangle, \sigma_1, \sigma_2) \vdash_\kappa \kappa : \tau';\ \psi^+ \mid \psi^-;\ o'\end{array}}{\tau_1;\ \psi^+_1 \mid \psi^-_1;\ o_1;\ \sigma_1 \vdash_\kappa \left\langle \widehat{\texttt{if}}_1\colon \widehat{v}, \rho, e_2, \rho_2, \sigma_2, \kappa\right\rangle^\circ : \tau';\ \psi^+ \mid \psi^-;\ o'}$$

$$\text{TK-SymIf2}^\circ\quad \frac{\begin{array}{c}\mathcal{E}(\rho_1,\sigma_1) \Vdash^\circ e_1 \ :\ \tau_2;\ \psi^+_1 \mid \psi^-_1; o_2 \\ \widehat{\tau}_2; \psi^+_2 \vee \psi^+_1 \mid \psi^-_2 \vee \psi^-_1; o_2; \mu_\sigma(\langle \widehat{v}, \rho\rangle, \sigma_1, \sigma_2) \vdash_\kappa \kappa : \tau';\ \psi^+ \mid \psi^-;\ o'\end{array}}{\tau_2;\ \psi^+_2 \mid \psi^-_2;\ o_2;\ \sigma_2 \vdash_\kappa \left\langle \widehat{\texttt{if}}_2\colon \widehat{v}, \rho, e_1, \rho_1, \sigma_1, \kappa\right\rangle^\circ : \tau';\ \psi^+ \mid \psi^-;\ o'}$$

$$\text{TK-App1}^\pi\quad \frac{\begin{array}{c}\mathcal{E}(\rho,\sigma) \Vdash^\pi e \ :\ \tau_1;\ \psi^+_e \mid \psi^-_e;\ o_e \\ \tau_2[x := o_e]; \psi^+[x := o_e] \mid \psi^-[x := o_e]; o[x := o_e]; \sigma \vdash_\kappa \kappa : \tau_3;\ \psi^+_3 \mid \psi^-_3;\ o_3\end{array}}{\pi; x\colon\tau_1 \xrightarrow{\psi^+ \mid \psi^-; o} \tau_2; \psi^+_{in} \mid \psi^-_{in}; o_{in}; \sigma \vdash_\kappa \langle \texttt{arg}\colon e, \rho, \kappa\rangle^\pi : \tau_3;\ \psi^+_3 \mid \psi^-_3;\ o_3}$$

$$\text{TK-App2}^\pi\quad \frac{\begin{array}{c}\mathcal{E}(\rho,\sigma) \Vdash^\pi v \ :\ \pi; x\colon\tau_1 \xrightarrow{\psi^+ \mid \psi^-; o} \tau_2;\ \psi^+_v \mid \psi^-_v;\ o_v \\ \tau_2[x := o_{in}]; \psi^+[x := o_{in}] \mid \psi^-[x := o_{in}]; o[x := o_{in}]; \sigma \vdash_\kappa \kappa : \tau_3;\ \psi^+_3 \mid \psi^-_3;\ o_3\end{array}}{\tau_1; \psi^+_{in} \mid \psi^-_{in}; o_{in}; \sigma \vdash_\kappa \langle \texttt{fn}\colon v, \rho, \kappa\rangle^\pi : \tau_3;\ \psi^+_3 \mid \psi^-_3;\ o_3}$$

$$\text{TK-Seq}^\pi\quad \frac{\mathcal{E}(\rho,\sigma) \Vdash^\pi e \ :\ \tau_e;\ \psi^+_e \mid \psi^-_e;\ o_e \qquad \tau_e;\ \psi^+_e \mid \psi^-_e;\ o_e; \sigma \vdash_\kappa \kappa : \tau;\ \psi^+ \mid \psi^-;\ o}{\tau_{in};\ \psi^+_{in} \mid \psi^-_{in};\ o_{in};\ \sigma \vdash_\kappa \langle \texttt{seq}\colon e, \rho, \kappa\rangle^\pi : \tau;\ \psi^+ \mid \psi^-;\ o}$$

Fig. 22. Type judgements for stack frames, part 1.

TK-Set!$^\bullet$

$$\frac{\mathcal{E}(\rho,\sigma) \overset{\bullet}{\vdash} v \ :\ \tau_{in};\ \psi_v^+ \mid \psi_v^-;\ o_v \qquad \overset{\textit{concrete?}\,\tau_{in}}{\text{where } \sigma[\ell] = \langle v, \rho\rangle} \qquad \texttt{Unit};\top;\bot;\cdot;\sigma \vdash_\kappa \kappa : \tau;\ \psi^+ \mid \psi^-;\ o}{\tau_{in};\ \psi_{in}^+ \mid \psi_{in}^-;\ o_{in};\ \sigma \vdash_\kappa \langle \texttt{set!}\colon \ell, \kappa\rangle^\bullet : \tau;\ \psi^+ \mid \psi^-;\ o}$$

TK-Set!-Sym$^\bullet$

$$\frac{\mathcal{E}(\rho,\sigma) \overset{\bullet}{\vdash} \widehat{v} \ :\ \widehat{\tau}_{in};\ \psi_v^+ \mid \psi_v^-;\ o_v \qquad \text{where } \sigma[\ell] = \langle \widehat{v}, \rho\rangle \qquad \texttt{Unit};\top;\bot;\cdot;\sigma \vdash_\kappa \kappa : \tau;\ \psi^+ \mid \psi^-;\ o}{\widehat{\tau}_{in};\psi_{in}^+ \mid \psi_{in}^-;o_{in};\sigma \vdash_\kappa \langle \texttt{set!}\colon \ell, \kappa\rangle^\bullet : \tau;\ \psi^+ \mid \psi^-;\ o}$$

TK-Set!$^\circ$

$$\frac{\mathcal{E}(\rho,\sigma) \overset{\circ}{\vdash} v \ :\ \tau_{in};\ \psi_v^+ \mid \psi_v^-;o_v \qquad \text{where } \sigma[\ell] = \langle \widehat{v}, \rho\rangle \qquad \texttt{Unit};\top;\bot;\cdot;\sigma \vdash_\kappa \kappa : \tau;\ \psi^+ \mid \psi^-;\ o}{\tau_{in};\psi_{in}^+ \mid \psi_{in}^-;o_{in};\sigma \vdash_\kappa \langle \texttt{set!}\colon \ell, \kappa\rangle^\circ : \tau;\ \psi^+ \mid \psi^-;\ o}$$

Fig. 23. Type judgements for stack frames, part 2: `set!` frames.

The $\vdash_\kappa$ rules in figure 22 are mostly straightforward, mirroring their counterparts in figure 13. Specifically, the rules for `if` use a symbolic $\overset{\circ}{\vdash}$ if the conditional test is symbolic, and the application rules require a function type with path concreteness that matches the marker on the stack frame. The interesting $\vdash_\kappa$ rules are for `set!` frames, in figure 23. When in a concrete path, TK-Set!$^\bullet$ and TK-Set!-Sym$^\bullet$ require the concreteness of τ_{in} to match the concreteness of the value already at location ℓ. When the path is symbolic, TK-Set!$^\circ$ specifies that type checking only succeeds if the value at location ℓ is symbolic.

With these definitions, we can state our key lemmas 4.3 and 4.4

Lemma 4.3 (Progress). *If* $\vdash_M M : \tau;\ \psi^+ \mid \psi^-;\ o$ *then either* $M \in VMach$ *or* $\exists M'$ *s.t.* $M \mapsto M'$.

Lemma 4.4 (Preservation). *If* $\vdash_M M : \tau;\ \psi^+ \mid \psi^-;\ o$ *and* $M \mapsto M'$, *then* $\vdash_M M' : \tau;\ \psi^{+\prime} \mid \psi^{-\prime};\ o'$.

Intuitively, lemma 4.3 states that well-typed terms are either values, or may make an evaluation step. The proof consists of a case analysis of the various possible combinations of control expressions and topmost stack frames. Lemma 4.4 states that machine transitions preserve well-typedness. The proof again proceeds by a case analysis of each machine transition rule. Together, the two lemmas are sufficient for proving theorem 4.1, and thus we may conclude that symbolic values cannot cause evaluation to get stuck.

5 TYPED ROSETTE

To show that $\widehat{\lambda}_\bullet$ can model a real programming language, we implemented Typed Rosette, which adds our type system to Torlak and Bodik (2014)'s untyped Rosette language.

5.1 Implementation Organization

Figure 24 depicts the overall architecture of Rosette and Typed Rosette, which are implemented with Racket (Flatt and PLT 2010). Racket's distinguishing features are its modern macro system (Flatt 2002, 2016)—a descendant of its Lisp and Scheme predecessors—and DSL-building capabilities (Tobin-Hochstadt et al. 2011). Together these features enable programmers to extend and reuse parts of Racket's implementation in order to create embedded DSLs quickly and easily. Rosette utilizes Racket's macro system to add symbolic computation capabilities and an interface to a solver. Since it is embedded in Racket, Rosette inherits Racket's DSL-building features and thus programmers may use it to easily build their own solver-aided DSLs (SDSLs).

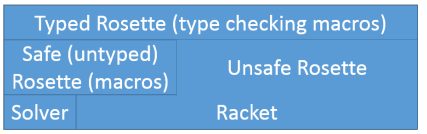

Fig. 24. Typed Rosette Organization

Rosette consists of two sublanguages: a small, safe subset where all language constructs are equipped to handle symbolic values, and a larger language that includes the rest of Racket. The documentation warns that programmers should stay within the safe subset, otherwise they must manually check correct program behavior and deal with problems emerging from mixing symbolic values and unlifted constructs. As with all languages, however, programmers are eventually enticed by the extra features and added expressiveness offered by the "unsafe" parts of the language. Typed Rosette supports all of unsafe Rosette language and thus aims to help the programmers that will inevitably want to use the unsafe language.

5.2 Implementation

Typed Rosette is implemented as an extension of Rosette, also using Racket's macro system. Creating Typed Rosette required implementing a type system, as well as an elaboration pass that translates the typed language to untyped Rosette. Since the elaboration pass occasionally requires type information to direct the elaboration output, e.g., when translating symbolic values like $\widehat{x}^{\mathtt{Int}}$ to its untyped representation, we did not implement separate passes. Instead, we used the "type systems as macros" technique of Chang et al. (2017), which allows implementing interleaved type checking and elaboration passes as a series of user-level macro definitions. With this approach is it trivial to implement elaborations that depend on type checking information. It also allowed us to use the existing Rosette implementation without modification. Overall Typed Rosette's implementation added roughly 2500 lines of code to the existing 10,000 lines of Rosette's implementation.

Figure 25 presents the essence of a few type checking macros from Typed Rosette's implementation. They use the `define-typechecking-macro` form from Chang et al.'s meta-DSL for creating typed DSLs. The `if` implementation on the left consists of two clauses, one each for concrete and symbolic test expressions, respectively, just like the rules from figure 8. The bodies of these `if` clauses also resemble their mathematical counterparts because its syntax is primarily a bidirectional (Pierce and Turner 1998) judgement that interleaves type checking and elaboration. Specifically, a judgement $\vdash e \gg \overline{e}\ (\Rightarrow k\ v) \ldots$ should be read "e elaborates to $\overline{e}$ and produces values $v \ldots$ associated with keys $k \ldots$." For example, the first premise in the first `if` clause elaborates the test expression e_1 to $\overline{e}_1$, and in the process computes its type τ, keyed with symbol `:`, as well as propositions ψ^+ and ψ^-. The next line guards the clause with a check of whether the test expression's type is symbolic. If the check fails, type checking falls through to the second clause. The next two premises type check the branches using the relevant propositions, and also toggle a flag to indicate a symbolic path. Finally, the output expression and its associated type lie below the conclusion line. Specifically, Typed Rosette's `if` elaborates to untyped Rosette's `if`. The output type is a join of the types of the two branches, computed with $\sqcup$. The symbolic clause turns its output type into a symbolic type.

The implementation for `set!` (figure 25, right) also has two clauses, for the cases of symbolic and concrete paths, respectively. When the path is symbolic, the type for the variable must also be symbolic. The final premise uses the $\Leftarrow$ arrow to "check" the type of e against the variable's type.

5.3 Concreteness Polymorphism in Practice

Typed Rosette extends $\widehat{\lambda}_{\mathfrak{d}}$ with additional features found in practical languages such as optional arguments and variable-arity polymorphism (Strickland et al. 2009). Unlike $\widehat{\lambda}_{\mathfrak{d}}$, where lambdas are

```
(define-typechecking-macro if
  [(if e₁ e₂ e₃) ≫ ; concrete case
    [⊢ e₁ ≫ ē₁ (⇒ : τ) (⇒ ψ⁺ ψ⁺) (⇒ ψ⁻ ψ⁻)]
    #:unless (and (symbolic? τ) (typecheck? τ False))
    [⊢ (with-prop e₂ ψ⁺) ≫ ē₂ ⇒ τ₁]
    [⊢ (with-prop e₃ ψ⁻) ≫ ē₃ ⇒ τ₂]
    ----------------------------------------
    [⊢ (rosette:if ē₁ ē₂ ē₃) ⇒ (⊔ τ₁ τ₂)]]
  [(if e₁ e₂ e₃) ≫ ; symbolic case
    [⊢ e₁ ≫ ē₁ (⇒ : τ) (⇒ ψ⁺ ψ⁺) (⇒ ψ⁻ ψ⁻)]
    [⊢ (sym-path (with-prop e₂ ψ⁺)) ≫ ē₂ ⇒ τ₁]
    [⊢ (sym-path (with-prop e₃ ψ⁻)) ≫ ē₃ ⇒ τ₂]
    ----------------------------------------
    [⊢ (rosette:if ē₁ ē₂ ē₃) ⇒ (⊔ τ₁ τ₂)]])
```

```
(define-typechecking-macro set!
  [(set! x e) ≫ #:when (sym-path?)
    [⊢ x ≫ x̄ ⇒ τ]
    #:fail-unless (symbolic? τ)
      "sym path requires sym type"
    [⊢ e ≫ ē ⇐ τ]
    ------------------------------
    [⊢ (rosette:set! x̄ ē) ⇒ Unit])
  [(set! x e) ≫ #:when (conc-path?)
    [⊢ x ≫ x̄ ⇒ τ]
    [⊢ e ≫ ē ⇐ τ]
    ------------------------------
    [⊢ (rosette:set! x̄ ē) ⇒ Unit])
```

Fig. 25. A few typechecking macros from Typed Rosette's implementation

always assigned types with all combinations of concrete/symbolic arguments, Typed Rosette programmers control which combinations are included in a function's type using signature declarations. For example, the concreteness of this function's output matches its first input:

```
#lang typed/rosette
(: add : (case-> (-> Int Int Int)
                 (-> Int Int Int)))
(define (add x y) (+ x y))
```

The `case->` constructor resembles the intersection type constructor from $\widehat{\lambda_{\mathfrak{d}}}$, except each constituent function type is considered in the listed order when the function is applied. The signature additionally specifies that the second input must always be concrete and does not affect the output type. Here is another example, where the second argument is now optional:

```
(: add/opt : (case-> (->* [Int] [Int] Int)
                     (->* [Int] [Int] Int)))
(define (add/opt x [y 0]) (+ x y))
```

The `->*` function type constructor requires three arguments: a list of types for required arguments, a list of types for optional arguments, and an output type. In both these cases, functions are implicitly assigned types considering both a symbolic and concrete path, as in $\widehat{\lambda_{\mathfrak{d}}}$.

5.4 Handling Imprecision

Concreteness polymorphism increases Typed Rosette's ability to track the concreteness of values. In some cases, however, Typed Rosette still rejects valid programs. Specifically, the merging function, e.g., $\mu_{\widehat{v}}$ from figure 18, gives rise to a discrepancy between the type system and the actual program. While the merging function for $\widehat{\lambda_{\mathfrak{d}}}$ always creates a symbolic value, an actual language should more aggressively preserve concrete values. For example, a value such as $\langle[\widehat{x}^{\texttt{Bool}}\colon 1][\neg\widehat{x}^{\texttt{Bool}}\colon 1]\rangle$ could be replaced with just the number 1. One of Rosette's key innovations is an novel merging algorithm that reduces the number of symbolic values during evaluation. For example, Typed Rosette rejects the following program, which uses the `add` function defined above:

```
(define-symbolic b : Bool)
(add 1 (if b 2 2)) ; => TYERR: expected Int given Int
```

The function does not type check because add's type requires a concrete second argument but the type system cannot prove this. To satisfy the type checker, we must use occurrence typing:

```
(let ([x (if b 2 2)])
  (if (conc? x)
      (add 1 x) ; => 3
      (error "expected concrete val")))
```

Another source of imprecision comes from Rosette's pruning of infeasible paths. For example, the following program on the left runs in untyped Rosette:

```
#lang rosette
(define-symbolic b boolean?)
(+ 1 (if b 2 "bad")) ; => 3
```

```
#lang typed/rosette
(define-symbolic b boolean?)
(+ 1 (assert-type (if b 2 "bad") : Int))
; => 3, with assertion b = true
```

Typed Rosette, however, rejects the program because the second argument's union type is incompatible with addition. To satisfy the type checker, one can use occurrence typing, like in the previous example, to restrict the type. Alternatively, Typed Rosette provides `assert-type` as seen on the right, which restricts the type of a value, but generates an assertion that the solver must satisfy.

6 EVALUATION

To determine whether Typed Rosette is useful, we ported a large code base from Rosette to Typed Rosette, summarized in table 1. We sought to confirm that Typed Rosette (1) sufficiently accommodates Rosette idioms and (2) helps with debugging lenient symbolic execution. Our test suite consists of two parts. The first part, summarized in the first line of table 1, consists of small examples mainly drawn from Rosette's documentation. The second part, summarized in the rest of the table, consists of example SDSLs.

The first part amounted to approximately 2000 lines of "coverage tests." While most used only safe Rosette, a few demonstrated errors from naive use of unsafe constructs. Typed Rosette was able to catch all these latter cases. The last column of table 1 shows that we needed 15 annotations to successfully type check these small test cases. All of these involved Rosette's pruning of infeasible paths as described in the section 5.4 and required adding `assert-type` annotations to satisfy the type checker. Since the documentation included a lot of corner case examples, however, it's not too surprising that we needed some annotations to help the type checker.

The second part of our test suite evaluated whether Typed Rosette is useful in practice. Specifically, we ported a series of SDSLs from Rosette to Typed Rosette, along with their accompanying test suites. Table 1 summarizes these efforts: the third column lists the size of the untyped code base; the fourth column shows how many lines were needed to add types; the fifth column shows how many lines of tests we ported to the typed version; and the last column shows how many type annotations we needed, if any. The rest of this section highlights a few cases in more detail.

6.1 Synthesizing Loop-Free Bitvector Programs

This SDSL provides forms for synthesizing bitvector programs in the manner described in Gulwani et al. (2011). The language is technically unsafe and Typed Rosette helps catch mistakes where symbolic values are used when concrete values are expected:

```
#lang typed/bv
(bv 1 4) ; constructs length 4 bitvector with value 1
(define-symbolic x : Int)
(bv x 4) ; => TYERR: expected Int, got x with type Int
```

Name	Description	Untyped LoC	+Typed LoC	Typed Tests	Anno.
tests	coverage tests, corner cases			1884	15
fsm	debugging automata	162	+86	54	
bv	synth loop-free bitvector progs	434	+101	438	
ifc	verify non-interference	962	+137	517	
synthcl	synth/verify opencl progs	2632	+615	1609	
ocelot	relational logic library	1511	+1954	934	8
inc	synth incremental algs	5445	+634	1134	12

Table 1. Summary of programs ported to Typed Rosette

Since the implementation is more or less complete, however, we were unable to find any errors in the code due to misused symbolic values. This is to be expected when porting mature code, however, and in the future we plan to use Typed Rosette for new projects where its debugging improvements may help more. This small language does illustrate how Rosette programmers usually manage interaction of symbolic values and unlifted constructed, however, by constructing additional abstractions on top of Rosette. For example, the following code tries to synthesize a function named `mk-trailing-0s-mask/synth` from a list of components and a reference implementation:

```
(define-fragment (mk-trailing-0s-mask/synth x)
  #:library (bvlib [(bv 1 4) bvsub bvand bvnot 1])
  #:implements mk-trailing-0s-mask)
```

The implementation of `define-fragment` performs synthesis by encoding the given components as symbolic values representing holes in a static single-assignment program, and then attempting to compute a satisfying assignment for those holes. Due to this abstraction, however, the user never interacts with the generated symbolic values, thus reducing the chance of making errors. Many of the other examples such as fsm, ifc, and synthcl are implemented in a similar manner.

6.2 A Library for Relational Logic Specifications

Despite the use of syntactic abstraction to minimize unsafe interactions, symbolic values can still occasionally leak through. This section reports on an example we encountered while porting the Ocelot library,[9] which extends Rosette with the ability to write, verify, and synthesize Alloy-like (Jackson 2002) relational specifications. Consider this example from the Ocelot documentation:

```
#lang rosette
(define Un (universe '(a b c d))) ; declare universe of atoms
(define cats (declare-relation 1 "cats")) ; declare a ”cats” relation
(define iCats (instantiate-bounds ; create an interpretation for ”cats”
                (bounds Un (list (make-upper-bound cats '((a) (b) (c) (d)))))))
(define F (and (some cats) (some (- univ cats)))) ; find an interesting model for ”cats”
(define resultCats (solve (assert (interpret* F iCats))))
; Lift the model back to atoms in Un
(interpretation->relations (evaluate iCats resultCats)) ; => cats: b
```

The code declares a universe of atoms `Un`; declares a relation `cats`; declares bounds for possible members of the relation; and compiles the relation to Rosette symbolic values using `instantiate-bounds`. Once the relation is compiled to base Rosette, the programmer may use Rosette's solver API, e.g., `solve`, to find a satisfying assignment of the symbolic values. Finally, the example uses

[9] https://jamesbornholt.github.io/ocelot/

`evaluate` to replace symbolic values with the concrete values computed by the solver, and then uses `interpretation->relations` to lift the encoded result back to the universe of "cats".

The `interpretation->relations` function utilizes unlifted code, however, and thus calling the function with a symbolic interpretation value is a mistake:

```
(interpretation->relations iCats) ; => cats: a,b,c,d (WRONG)
```

Despite the mistake the program does not produce an error; instead, it silently returns the wrong answer because the symbolic values are mistakenly interpreted as "true" values. With our typed version, the erroneous code above produces an error:

```
#lang typed/rosette
(interpretation->relations iCats) ; => TYERR: expected concrete value
```

Other parts of Ocelot exposed a gap in our type system involving Racket `structs`, which are a kind of named record. A struct definition may "inherit" properties from a base struct and thus type checking them requires record subtyping. Typed Rosette's occurrence typing, however, does not currently support refinement of this kind of record subtyping and thus we needed casts and dynamic checks in a few places to fully type check this library.

6.3 Synthesizing Incremental Algorithms

Incremental algorithms speed up programs by avoiding full recomputation when successive inputs are related. We ported an SDSL for synthesizing incremental algorithms (Shah and Bodik 2017)[10] from Rosette to Typed Rosette. The language utilizes the unsafe Rosette variant, in particular hash tables, and relies on dynamic checks to prevent symbolic values from reaching unlifted constructs.

By porting the code to Typed Rosette, we were able to remove approximately two dozen of these manually-inserted checks (approximately 100 lines). Of these two dozen, we discovered (and the authors acknowledged) that about half were erroneous or incomplete, meaning it was still possible for symbolic values to crash the program by flowing to locations that could not handle them. Our typed version was able to catch and prevent all these errors.

We needed two kinds of annotations to make this library work with Typed Rosette. The first involved either lifting a type to be symbolic for mutation purposes. For example:

```
#lang typed/rosette
(define v (make-vector 7 (ann #f : (U Bool))))
(define-symbolic b : Bool)
(if b (vector-set! v 0 #t) (vector-set! v 1 #t))
```

This code creates a 7-element vector initialized with concrete false values and would thus be assigned a concrete type. We wanted to mutate the vector in a symbolic path, however, so we needed to annotate the initialization value with a symbolic type.

The second class of annotations we needed involved hash tables. Specifically, the result of a hash lookup has a union type because the lookup may return false the lookup fails. Thus to use the result of the lookup, we need an extra occurrence typing check to eliminate the false case:

```
#lang typed/rosette
(define h (hash 'a 0 'b 1))
(let ([x (hash-ref h 'a #f)]) ; x has type (U Int False) due to fail case
  (if (false? x) (error "failed!") (+ x 1))) ; x has type Int
```

[10] Thanks to Rohin Shah for providing us access to the code repository.

7 RELATED WORK

Handling Unlifted Constructs Most symbolic execution engines prevent invalid interaction of symbolic values with unlifted constructs by either equipping an entire language to handle symbolic values or limiting programmers to a safe subset. For example, Symbolic Pathfinder (Păsăreanu et al. 2008) symbolically executes Java programs with a replacement JVM supporting symbolic values. This approach is the more difficult, however, and may produce complex solver encodings. Alternatively, the Leon (Blanc et al. 2013) and Kaplan (Köksal et al. 2012) languages restrict programmers to a subset of Scala called PureScala. Similarly, Rubicon (Near and Jackson 2012) symbolically executes only a subset of Ruby but does not support interacting with other parts of the language. This approach limits the expressiveness of the language and thus possible applications of the tool.

The above approaches of ensuring safe behavior of symbolic values may still be insufficient if the program interacts with components, like libraries, outside the control of the symbolic execution engine. Concolic testing frameworks in particular have devised various heuristics to handle this situation. The EXE (Cadar et al. 2006) framework, a tool for testing C programs, handles uninstrumented parts of the program by logging calls into uninstrumented functions, instrumenting the code, rerunning the program, and then repeating the process until there are no more uninstrumented function calls. It's not clear, however, if this algorithm accommodates dynamically-linked libraries. Other tools, like CUTE (Sen et al. 2005) and KLEE (Cadar et al. 2008) concretize symbolic values when they flow to library code. This approach does not consider all program paths and thus may not work for some applications of symbolic execution such as verification. The DART (Godefroid et al. 2005) tool concretizes symbolic values but trades incompleteness for possible non-termination. Specifically, when encountering uninstrumented code, execution falls back to the concrete result and the tool sets an "incompleteness" flag to true. Testing continues until all branches are covered; otherwise the tool runs forever.

SMTen resembles Rosette in that it allows specifying search problems with a high-level functional language, Haskell, which is then compiled to solver encodings. Programmers do not explicitly compute with symbolic values, however, making SMTen less general than Rosette. Thus, although it is a typed language, SMTen does not utilize symbolic types in the manner of Typed Rosette.

Tracking Value Flow Our type system resembles a static taint analysis, which is commonly used to prove security properties such as non-interference. These analyses, however, are typically complex, sometimes even more complicated than symbolic execution itself since they must track the flow of individual values. Also, taint analyses are typically computed at the assembly or bytecode level, which may pose challenges when trying to use it for debugging high-level programs. In contrast, our system works with high-level programs, and we intentionally sought an algorithm that is lightweight enough to be considered "type checking", yet still precise enough to be useful.

The **Mix** system (Khoo et al. 2010) combines type checking and symbolic execution but utilizes them as alternatives. Specifically, it applies one or the other to different parts of a program, gaining efficiency over symbolic execution alone. Its type system, however, does not distinguish symbolic values from non-symbolic ones and execution does not mix symbolic values with unlifted constructs.

8 CONCLUSION

We advocate for "lenient symbolic execution", which potentially is easier to implement, produces simpler encodings, and allows programmers to benefit from a full range of language features. To ensure safe behavior of symbolic values, a type checker should accompany such a mixed system. To this end, we present $\widehat{\lambda}_{\mathfrak{d}}$, a typed λ-calculus whose types distinguish symbolic from concrete values. Our calculus is useful enough to model Typed Rosette, a typed solver-aided language supporting lenient symbolic execution, and we evaluated Typed Rosette with a comprehensive test suite.

REFERENCES

Régis Blanc, Viktor Kuncak, Etienne Kneuss, and Philippe Suter. 2013. An Overview of the Leon Verification System: Verification by Translation to Recursive Functions. In *Proceedings of the 4th Workshop on Scala*. 1:1–1:10.

James Bornholt and Emina Torlak. 2017. Synthesizing Memory Models from Framework Sketches and Litmus Tests. In *Proceedings of the 38th ACM SIGPLAN Conference on Programming Language Design and Implementation (PLDI 2017)*. 467–481.

James Bornholt, Emina Torlak, Dan Grossman, and Luis Ceze. 2016. Optimizing Synthesis with Metasketches. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. 775–788.

Robert S. Boyer, Bernard Elspas, and Karl N. Levitt. 1975. SELECT—a Formal System for Testing and Debugging Programs by Symbolic Execution. In *Proceedings of the International Conference on Reliable Software*. 234–245.

Cristian Cadar, Daniel Dunbar, and Dawson Engler. 2008. KLEE: Unassisted and Automatic Generation of High-coverage Tests for Complex Systems Programs. In *Proceedings of the 8th USENIX Conference on Operating Systems Design and Implementation*. 209–224.

Cristian Cadar, Vijay Ganesh, Peter M. Pawlowski, David L. Dill, and Dawson R. Engler. 2006. EXE: Automatically Generating Inputs of Death. In *Proceedings of the 13th ACM Conference on Computer and Communications Security*. 322–335.

Stephen Chang, Alex Knauth, and Ben Greenman. 2017. Type Systems As Macros. In *Proceedings of the 44th ACM SIGPLAN Symposium on Principles of Programming Languages*. 694–705.

Greg Dennis, Felix Sheng-Ho Chang, and Daniel Jackson. 2006. Modular Verification of Code with SAT. In *Proceedings of the 2006 International Symposium on Software Testing and Analysis*. 109–120.

Azadeh Farzan, Andreas Holzer, Niloofar Razavi, and Helmut Veith. 2013. Con2Colic Testing. In *Proceedings of the 2013 9th Joint Meeting on Foundations of Software Engineering*. 37–47.

Matthias Felleisen, Robert Bruce Findler, and Matthew Flatt. 2009. *Semantics Engineering with PLT Redex*. MIT Press.

Matthew Flatt. 2002. Composable and Compilable Macros: You Want It when?. In *Proceedings of the Seventh ACM SIGPLAN International Conference on Functional Programming*. 72–83.

Matthew Flatt. 2016. Binding As Sets of Scopes. In *Proceedings of the 43rd Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. 705–717.

Matthew Flatt and PLT. 2010. *Reference: Racket*. Technical Report PLT-TR-2010-1. PLT Design Inc. https://racket-lang.org/tr1/.

Tim Freeman and Frank Pfenning. 1991. Refinement Types for ML. In *Proceedings of the ACM SIGPLAN 1991 Conference on Programming Language Design and Implementation*. 268–277.

Patrice Godefroid, Nils Klarlund, and Koushik Sen. 2005. DART: Directed Automated Random Testing. In *Proceedings of the 2005 ACM SIGPLAN Conference on Programming Language Design and Implementation*. 213–223.

Patrice Godefroid, Michael Y. Levin, and David Molnar. 2012. SAGE: Whitebox Fuzzing for Security Testing. *Queue* 10, 1 (2012), 20:20–20:27.

Sumit Gulwani, Susmit Jha, Ashish Tiwari, and Ramarathnam Venkatesan. 2011. Synthesis of Loop-free Programs. In *Proceedings of the 32nd ACM SIGPLAN Conference on Programming Language Design and Implementation*. 62–73.

Daniel Jackson. 2002. Alloy: A Lightweight Object Modelling Notation. *ACM Trans. Softw. Eng. Methodol.* 11, 2 (2002), 256–290.

Joxan Jaffar, Vijayaraghavan Murali, Jorge A. Navas, and Andrew E. Santosa. 2012. TRACER: A Symbolic Execution Tool for Verification. In *Computer Aided Verification: 24th International Conference, CAV 2012, Berkeley, CA, USA, July 7-13, 2012 Proceedings*. 758–766.

Yit Phang Khoo, Bor-Yuh Evan Chang, and Jeffrey S. Foster. 2010. Mixing Type Checking and Symbolic Execution. In *Proceedings of the 31st ACM SIGPLAN Conference on Programming Language Design and Implementation*. 436–447.

James C. King. 1975. A New Approach to Program Testing. In *Proceedings of the International Conference on Reliable Software*. 228–233.

Ali Sinan Köksal, Viktor Kuncak, and Philippe Suter. 2012. Constraints As Control. In *Proceedings of the 39th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. 151–164.

Joseph P. Near and Daniel Jackson. 2012. Rubicon: Bounded Verification of Web Applications. In *Proceedings of the ACM SIGSOFT 20th International Symposium on the Foundations of Software Engineering*. 60:1–60:11.

Stuart Pernsteiner, Calvin Loncaric, Emina Torlak, Zachary Tatlock, Xi Wang, Michael D. Ernst, and Jonathan Jacky. 2016. Investigating Safety of a Radiotherapy Machine Using System Models with Pluggable Checkers. In *Computer Aided Verification: 28th International Conference, CAV 2016, Toronto, ON, Canada, July 17-23, 2016, Proceedings, Part II*. 23–41.

Phitchaya Mangpo Phothilimthana, Tikhon Jelvis, Rohin Shah, Nishant Totla, Sarah Chasins, and Rastislav Bodik. 2014. Chlorophyll: Synthesis-aided Compiler for Low-power Spatial Architectures. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*. 396–407.

Benjamin C. Pierce and David N. Turner. 1998. Local Type Inference. In *Proceedings of the 25th ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages*. 252–265.

Corina S. Păsăreanu, Peter C. Mehlitz, David H. Bushnell, Karen Gundy-Burlet, Michael Lowry, Suzette Person, and Mark Pape. 2008. Combining Unit-level Symbolic Execution and System-level Concrete Execution for Testing NASA Software. In *Proceedings of the 2008 International Symposium on Software Testing and Analysis*. ACM, New York, NY, USA, 15–26.

Koushik Sen, Darko Marinov, and Gul Agha. 2005. CUTE: A Concolic Unit Testing Engine for C. In *Proceedings of the 10th European Software Engineering Conference Held Jointly with 13th ACM SIGSOFT International Symposium on Foundations of Software Engineering*. 263–272.

Rohin Shah and Rastislav Bodik. 2017. Automated Incrementalization through Synthesis. In *Proceedings of the First Workshop on Incremental Computing*.

Armando Solar-Lezama, Rodric Rabbah, Rastislav Bodík, and Kemal Ebcioğlu. 2005. Programming by Sketching for Bit-streaming Programs. In *Proceedings of the 2005 ACM SIGPLAN Conference on Programming Language Design and Implementation*. 281–294.

T. Stephen Strickland, Sam Tobin-Hochstadt, and Matthias Felleisen. 2009. Practical Variable-Arity Polymorphism. In *Proceedings of the 18th European Symposium on Programming Languages and Systems: Held As Part of the Joint European Conferences on Theory and Practice of Software, ETAPS 2009*. 32–46.

Sam Tobin-Hochstadt and Matthias Felleisen. 2010. Logical Types for Untyped Languages. In *Proceedings of the 15th ACM SIGPLAN International Conference on Functional Programming*. 117–128.

Sam Tobin-Hochstadt, Vincent St-Amour, Ryan Culpepper, Matthew Flatt, and Matthias Felleisen. 2011. Languages As Libraries. In *Proceedings of the 32Nd ACM SIGPLAN Conference on Programming Language Design and Implementation*. 132–141.

Emina Torlak and Rastislav Bodik. 2013. Growing Solver-aided Languages with Rosette. In *Proceedings of the 2013 ACM International Symposium on New Ideas, New Paradigms, and Reflections on Programming & Software*. 135–152.

Emina Torlak and Rastislav Bodik. 2014. A Lightweight Symbolic Virtual Machine for Solver-Aided Host Languages. In *Proceedings of the 35th ACM SIGPLAN Conference on Programming Language Design and Implementation*. 530–541.

Konstantin Weitz, Doug Woos, Emina Torlak, Michael D. Ernst, Arvind Krishnamurthy, and Zachary Tatlock. 2016. Scalable Verification of Border Gateway Protocol Configurations with an SMT Solver. In *Proceedings of the 2016 ACM SIGPLAN International Conference on Object-Oriented Programming, Systems, Languages, and Applications*. 765–780.

A.K. Wright and M. Felleisen. 1994. A Syntactic Approach to Type Soundness. *Inf. Comput.* 115, 1 (1994), 38–94.